

4.3 Exercise: random intercepts and slopes

Benjamin Rosenbaum

November 5, 2021

We will extend the last model by including random slopes, too.

Setup

```
rm(list=ls())
library(rstan)
library(coda)
library(BayesianTools)
library(brms)

setwd("~/Nextcloud/teaching Bayes 2021")

rstan_options(auto_write = TRUE)
options(mc.cores = 4)
```

Read dataset


```
df = read.table("data/RIKZ.txt", header=TRUE)
head(df)

##   Sample Richness Exposure    NAP Beach
## 1      1      11      10  0.045     1
## 2      2      10      10 -1.036     1
## 3      3      13      10 -1.336     1
## 4      4      11      10  0.616     1
## 5      5      10      10 -0.684     1
## 6      6       8       8  1.190     2

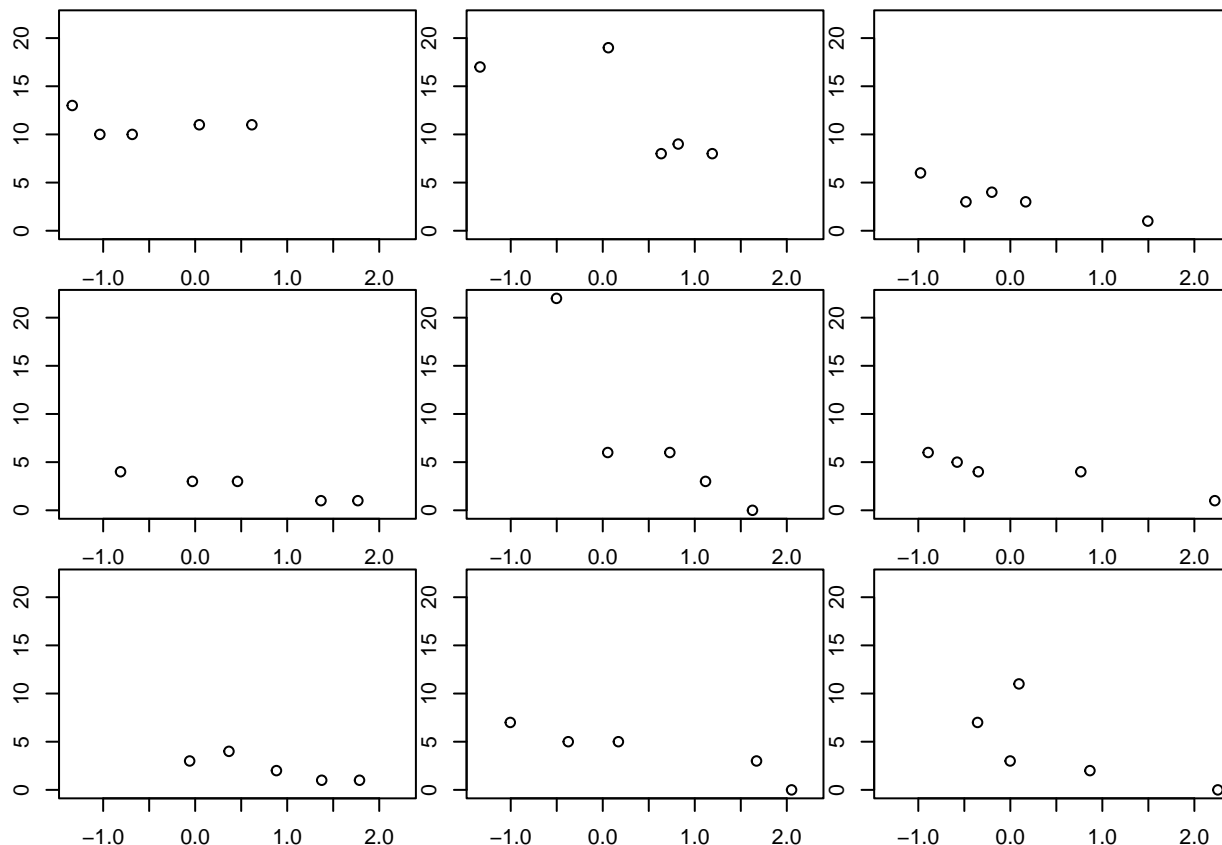
str(df)

## 'data.frame':   45 obs. of  5 variables:
##  $ Sample : int  1 2 3 4 5 6 7 8 9 10 ...
##  $ Richness: int  11 10 13 11 10 8 9 8 19 17 ...
##  $ Exposure: int  10 10 10 10 10 8 8 8 8 8 ...
##  $ NAP : num  0.045 -1.036 -1.336 0.616 -0.684 ...
##  $ Beach : int  1 1 1 1 1 2 2 2 2 2 ...

par(mfrow=c(1,1))
plot(df$NAP, df$Richness)
```



```
par(mfrow=c(3,3), mar=c(1,1,1,1), oma=c(1,1,0,0))
for (i in 1:9){
  df.sub=subset(df, df$Beach==i)
  plot(df.sub$NAP, df.sub$Richness,
       xlim=range(df$NAP),
       ylim=range(df$Richness) )
}
```



Random intercepts model

We will fit linear regression lines to each group (**Beach**) as follows:

$$\begin{aligned}
 y_i &\sim \text{normal}(a_{\text{group}(i)} + b_{\text{group}(i)} \cdot x_i, \sigma) \quad i = 1, \dots, n \quad (n \text{ observations}) \\
 a_j &\sim \text{normal}(\mu_a, \sigma_a) \quad j = 1, \dots, m \quad (m \text{ groups}) \\
 b_j &\sim \text{normal}(\mu_b, \sigma_b) \quad j = 1, \dots, m \quad (m \text{ groups})
 \end{aligned}$$

Here, a_j and b_j are group-level intercepts and slopes, which are allowed to vary (partial pooling).

Both have their own means (μ_a , μ_b) and standard deviations (σ_a , σ_b), which are also free parameters to be estimated.

So this is a random intercepts and slopes linear regression, lm-formulation would be $y \sim x + (x|\text{group})$, which is short for $y \sim 1+x + (1+x|\text{group})$.

Here, an intercept a_j and a slope b_j are independent. Usually they are modeled with a correlation (standard in `lme4` and also in `brms`), but we leave that out in Stan.

```
data = list(y = df$Richness,
            x = df$NAP,
            group = df$Beach,
            n = nrow(df),
            n_group = 9)
```

```
data
```

```
## $y
```

```
## [1] 11 10 13 11 10 8 9 8 19 17 6 1 4 3 3 1 3 3 1 4 3 22 6 0 6
## [26] 5 4 1 6 4 2 1 1 3 4 3 5 7 5 0 7 11 3 0 2
##
## $x
## [1] 0.045 -1.036 -1.336 0.616 -0.684 1.190 0.820 0.635 0.061 -1.334
## [11] -0.976 1.494 -0.201 -0.482 0.167 1.768 -0.030 0.460 1.367 -0.811
## [21] 1.117 -0.503 0.729 1.627 0.054 -0.578 -0.348 2.222 -0.893 0.766
## [31] 0.883 1.786 1.375 -0.060 0.367 1.671 -0.375 -1.005 0.170 2.052
## [41] -0.356 0.094 -0.002 2.255 0.865
##
## $group
## [1] 1 1 1 1 1 2 2 2 2 2 3 3 3 3 3 4 4 4 4 4 5 5 5 5 5 6 6 6 6 6 7 7 7 7 7 8 8 8
## [39] 8 8 9 9 9 9 9
##
## $n
## [1] 45
##
## $n_group
## [1] 9
```

```
stan_code_partpool = '
data {
  int n;
  int n_group;
  real y[n];
  real x[n];
  int group[n];
}
parameters {
  real a[n_group];
  real b[n_group];
  real<lower=0> sigma;
  real mu_a;
  real<lower=0> sigma_a;
  real mu_b;
  real<lower=0> sigma_b;
}
model {
  // priors
  mu_a ~ normal(0,10);
  mu_b ~ normal(0,10);
  sigma_a ~ cauchy(0,1);
  sigma_b ~ cauchy(0,1);
  for (j in 1:n_group){
    a[j] ~ normal(mu_a,sigma_a);
    b[j] ~ normal(mu_b,sigma_b);
  }
  sigma ~ normal(0,10);
  // likelihood
  for(i in 1:n){
    y[i] ~ normal( a[ group[i] ] + b[ group[i] ] *x[i] , sigma);
  }
}
'
```


```

stan_model_partpool = stan_model(model_code=stan_code_partpool)
fit_partpool = sampling(stan_model_partpool, data=data)


print(fit_partpool, digits=3, probs=c(0.025, 0.975))

## Inference for Stan model: 09c04efcb9dcbe2facbb0e83ef1aa41b.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##               mean se_mean    sd      2.5%    97.5% n_eff  Rhat
## a[1]          9.346   0.032 1.360    6.621   12.009  1826 1.000
## a[2]         11.861   0.048 1.493    8.528   14.473   956 1.002
## a[3]          4.016   0.032 1.283    1.606    6.608  1653 1.002
## a[4]          4.198   0.041 1.355    1.547    6.855  1089 1.004
## a[5]          9.307   0.048 1.598    6.281   12.505  1106 1.003
## a[6]          4.907   0.032 1.316    2.344    7.536  1736 1.001
## a[7]          4.762   0.031 1.434    1.933    7.574  2161 1.001
## a[8]          5.477   0.041 1.281    2.964    8.130   997 1.004
## a[9]          6.218   0.024 1.305    3.662    8.969  3046 1.001
## b[1]         -2.179   0.028 1.069   -4.032    0.241  1484 1.005
## b[2]         -2.991   0.024 0.978   -5.098   -1.166  1704 1.002
## b[3]         -2.380   0.020 0.933   -4.201   -0.437  2233 1.002
## b[4]         -2.272   0.025 0.930   -4.028   -0.286  1352 1.005
## b[5]         -4.144   0.063 1.637   -7.857   -1.759   679 1.003
## b[6]         -2.175   0.024 0.851   -3.763   -0.369  1227 1.004
## b[7]         -2.518   0.022 0.990   -4.495   -0.471  1971 1.003
## b[8]         -2.375   0.021 0.776   -3.891   -0.753  1341 1.004
## b[9]         -2.735   0.020 0.877   -4.556   -1.071  1860 1.003
## sigma         2.981   0.014 0.454    2.241    4.052  1097 1.005
## mu_a          6.594   0.024 1.191    4.199    9.010  2455 1.001
## sigma_a       3.169   0.034 1.103    1.417    5.767  1060 1.003
## mu_b         -2.626   0.017 0.639   -3.880   -1.426  1427 1.006
## sigma_b       1.052   0.033 0.706    0.187    2.733   472 1.004
## lp__        -89.426   0.483 6.067 -100.416  -74.555   158 1.010
##
## Samples were drawn using NUTS(diag_e) at Tue Oct 12 14:12:38 2021.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).
plot(fit_partpool, pars="a")

```



```
plot(fit_partpool, pars="b")
```



Predictions / credible intervals

Again, we can generate predictions and compute credible intervals (for the deterministic part of the model). Here: 90% credible intervals.

```

posterior=as.matrix(fit_partpool)

par(mfrow=c(3,3), mar=c(1,1,1,1), oma=c(1,1,0,0))
for (i in 1:9){
  df.sub=subset(df, df$Beach==i)
  x.pred = seq(from=min(df.sub$NAP), to=max(df.sub$NAP), by=0.01)
  plot(df.sub$NAP, df.sub$Richness,
       xlim=range(df$NAP),
       ylim=range(df$Richness) )

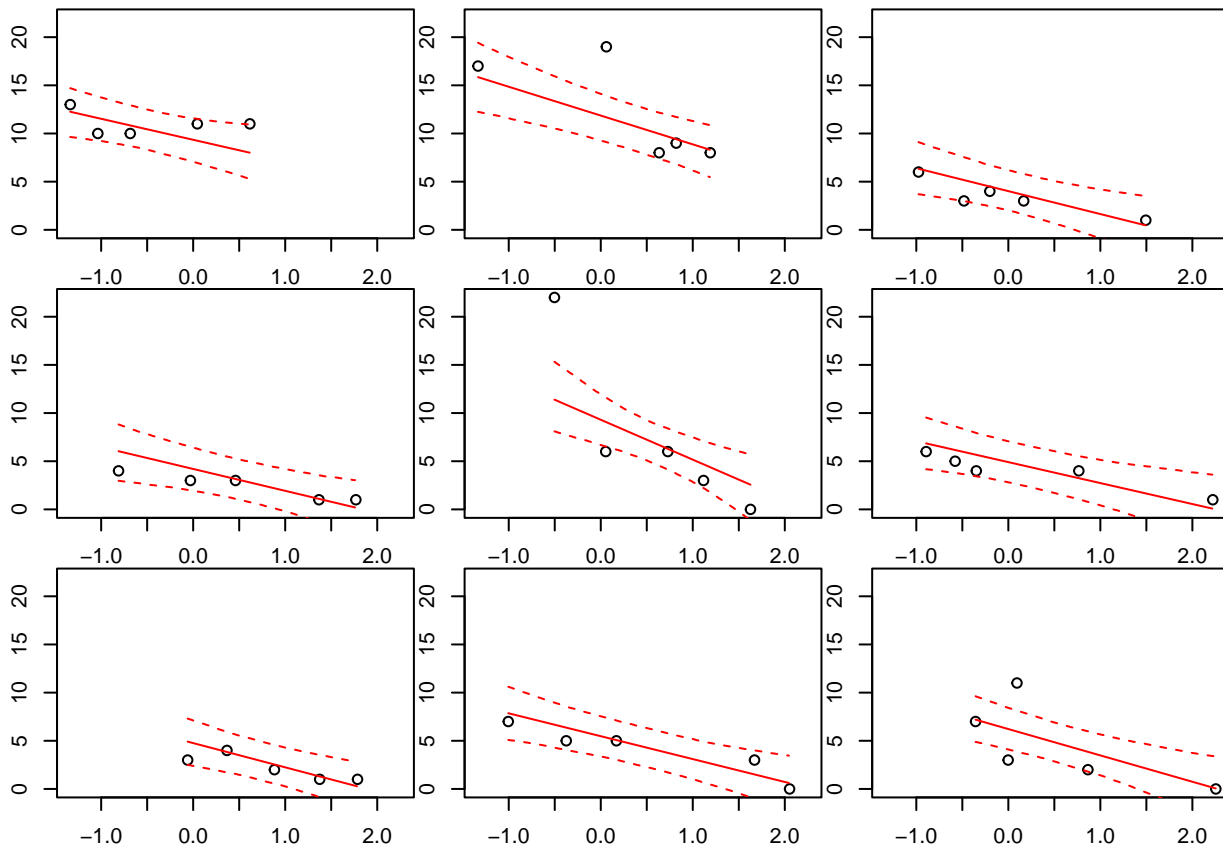
  y.cred = matrix(0, nrow=nrow(posterior), ncol=length(x.pred))
  for(j in 1:nrow(posterior)){
    y.cred[j, ] = posterior[j,paste0("a[",i,""])] + posterior[j,paste0("b[",i,""])]*x.pred
  }

  y.cred.mean = apply(y.cred, 2, function(x) mean(x))
  lines(x.pred, y.cred.mean, col="red")

  y.cred.q05 = apply(y.cred, 2, function(x) quantile(x, probs=0.05))
  lines(x.pred, y.cred.q05, col="red", lty=2)

  y.cred.q95 = apply(y.cred, 2, function(x) quantile(x, probs=0.95))
  lines(x.pred, y.cred.q95, col="red", lty=2)
}

```




brms complete pooling

We start with the “complete pooling” model in `brms`, which fits just a linear regression on `NAP`, ignoring the effect of predictor `Beach` on the intercept and slope. It’s the same complete pooling model as in the random intercepts regression before.

```
fit.b.compl = brm( Richness ~ NAP,  
                  data=df )
```

```
fit.b.compl
```

```
## Family: gaussian  
## Links: mu = identity; sigma = identity  
## Formula: Richness ~ NAP  
## Data: df (Number of observations: 45)  
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;  
## total post-warmup draws = 4000  
##  
## Population-Level Effects:  
##      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS  
## Intercept      6.64      0.67    5.31    7.94 1.00     3382     2541  
## NAP           -2.87      0.64   -4.13   -1.60 1.00     3834     2667  
##  
## Family Specific Parameters:  
##      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS  
## sigma      4.24      0.46    3.46    5.25 1.00     3193     2526  
##  
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS  
## and Tail_ESS are effective sample size measures, and Rhat is the potential  
## scale reduction factor on split chains (at convergence, Rhat = 1).  
  
plot( conditional_effects(fit.b.compl),  
      points=TRUE,  
      ask=FALSE )
```

brms no pooling

Next, “no pooling” fits different intercepts and slopes for all level of the categorical predictor **Beach**.

But first we must code **Beach** as a factor, otherwise it would be interpreted as a continuous predictor.

```
str(df)
```

```
## 'data.frame':   45 obs. of  5 variables:
##  $ Sample   : int   1  2  3  4  5  6  7  8  9 10 ...
##  $ Richness : int   11 10 13 11 10  8  9  8 19 17 ...
##  $ Exposure : int   10 10 10 10 10  8  8  8  8  8 ...
##  $ NAP      : num   0.045 -1.036 -1.336 0.616 -0.684 ...
##  $ Beach    : int    1  1  1  1  1  2  2  2  2  2 ...
```

```
df$Beach = as.factor(df$Beach)
```

```
fit.b.no = brm( Richness ~ NAP * Beach,
               data=df )
```

As usual with categorical predictors, `lm()` or `brm()` uses dummy coding. I.e. the presented effects are not the different intercepts and slopes for all levels of **Beach**. The intercept for the first level of **Beach** is **Intercept**, but **Beach2** etc are the differences in intercept for the other levels. **NAP** (effect name **b_NAP**) is the slope for the first level of **Beach**, and **NAP:Beach2** etc are the differences in slope for the other levels.


```
fit.b.no
```

```
## Family: gaussian
## Links: mu = identity; sigma = identity
## Formula: Richness ~ NAP * Beach
```

```
## Data: df (Number of observations: 45)
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
## total post-warmup draws = 4000
##
## Population-Level Effects:
##      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept      10.78      1.42      8.04     13.53 1.00      992     1576
## NAP            -0.39      1.60     -3.53      2.64 1.00      892     1480
## Beach2          2.56      1.85     -1.00      6.12 1.00     1422     2096
## Beach3         -7.38      1.81    -10.90     -3.82 1.00     1385     1888
## Beach4         -7.69      1.96    -11.55     -3.74 1.00     1498     2209
## Beach5          2.00      2.05     -2.07      6.02 1.00     1752     2229
## Beach6         -6.45      1.82    -10.08     -2.92 1.00     1363     2109
## Beach7         -7.22      2.44    -12.03     -2.36 1.00     1821     2228
## Beach8         -5.84      1.91     -9.60     -2.08 1.00     1376     1955
## Beach9         -4.52      1.96     -8.39     -0.55 1.00     1509     2229
## NAP:Beach2     -3.77      2.03     -7.73      0.24 1.00     1290     2174
## NAP:Beach3     -1.37      2.07     -5.34      2.60 1.00     1262     1998
## NAP:Beach4     -0.87      2.00     -4.74      3.11 1.00     1200     2164
## NAP:Beach5     -8.54      2.23    -12.95     -4.14 1.00     1295     2330
## NAP:Beach6     -1.00      1.87     -4.57      2.67 1.00     1118     1930
## NAP:Beach7     -1.17      2.38     -5.78      3.73 1.00     1462     1829
## NAP:Beach8     -1.52      1.87     -5.09      2.16 1.00     1201     1788
## NAP:Beach9     -2.57      2.03     -6.50      1.55 1.00     1321     2150
##
## Family Specific Parameters:
##      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma      2.57      0.37      1.96      3.42 1.00     1726     1993
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

By using conditions= and specifying all levels of Beach, effect of NAP is shown for all levels. Slope is different in all levels.

```
plot( conditional_effects(fit.b.no,
                          effects = "NAP",
                          conditions = data.frame( Beach=levels(df$Beach) ) ),
      points=TRUE,
      ask=FALSE )
```



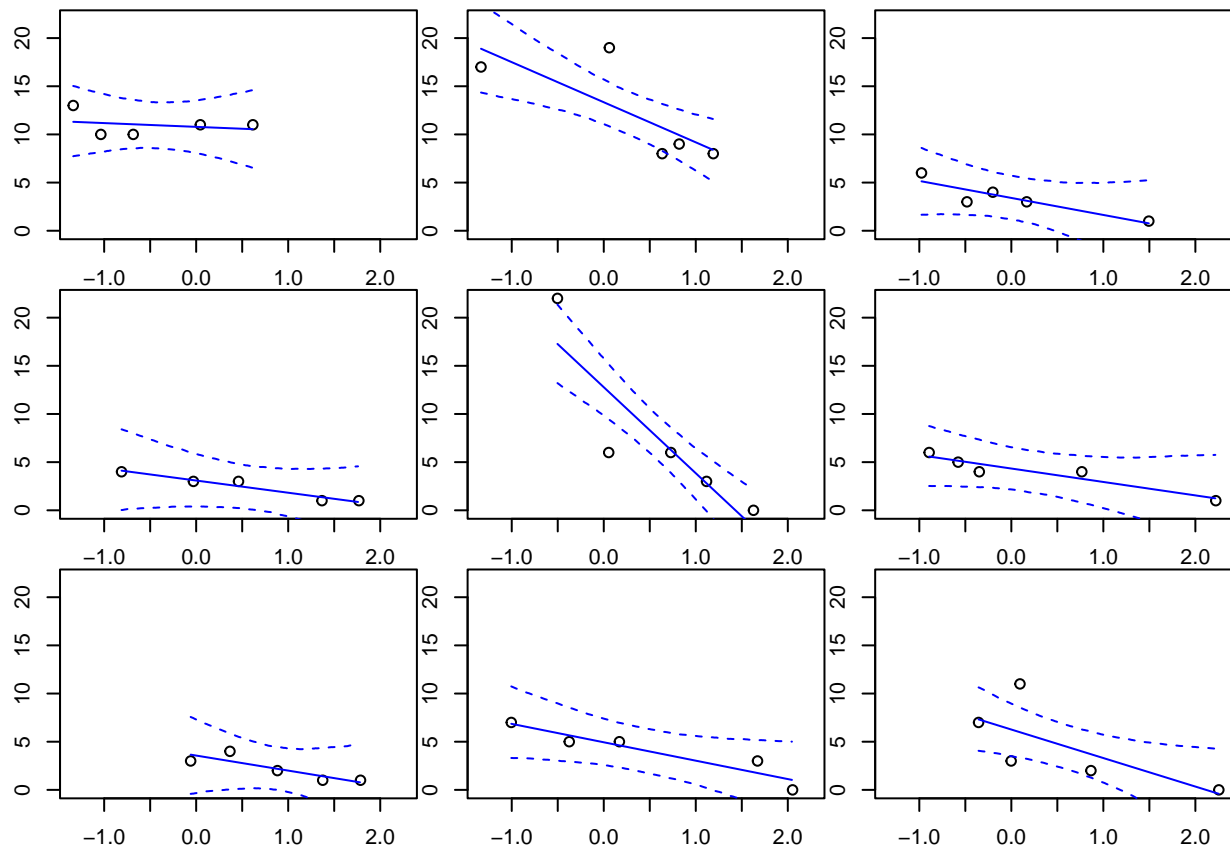
The group-level predictions can also be extracted by hand using `fitted()`. We can specify the range of NAP and levels of Beach with `newdata`. We plot 95% credible intervals.

```
levels = levels(df$Beach)
par(mfrow=c(3,3), mar=c(1,1,1,1), oma=c(1,1,0,0))

for (i in 1:9){
  df.sub=subset(df, df$Beach==i)
  x.pred = seq(from=min(df.sub$NAP), to=max(df.sub$NAP), by=0.01)
  plot(df.sub$NAP, df.sub$Richness,
       xlim=range(df$NAP),
       ylim=range(df$Richness) )

  y.cred = fitted(fit.b.no, newdata=data.frame(NAP=x.pred,
                                              Beach=levels[i] ) )


  lines(x.pred, y.cred[, 1], col="blue")
  lines(x.pred, y.cred[, 3], col="blue", lty=2)
  lines(x.pred, y.cred[, 4], col="blue", lty=2)
}
```



Or we can plot the whole dataset with all group-level mean regression lines.

```
plot(df$NAP, df$Richness)
```

```
for (i in 1:9){
  df.sub=subset(df, df$Beach==i)
  x.pred = seq(from=min(df.sub$NAP), to=max(df.sub$NAP), by=0.01)
  y.cred = fitted(fit.b.no, newdata=data.frame(NAP=x.pred,
                                                Beach=levels[i] ) )
  lines(x.pred, y.cred[, 1], col="blue")
}
```



brms partial pooling

With partial pooling, we use a random intercepts and slopes for the categorical predictor `Beach`. This time, we provide priors for the “fixed effects”, i.e. mean intercept and slope.

```
priors = c(prior(normal(5,5), class=Intercept),
           prior(normal(0,10), class=b) )
```

```
fit.b.part = brm( Richness ~ 1+NAP + (1+NAP|Beach),
                  prior=priors,
                  data=df )
```

Now, in the population-level effects mean intercept μ_a and slope μ_b are shown. The group-level effects only contain a summary of the random effects. It's standard deviations (σ_a, σ_b) and the correlation between slopes and intercepts $\text{cor}(a, b)$.

```
fit.b.part
```

```
## Family: gaussian
## Links: mu = identity; sigma = identity
## Formula: Richness ~ 1 + NAP + (1 + NAP | Beach)
## Data: df (Number of observations: 45)
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
## total post-warmup draws = 4000
##
## Group-Level Effects:
## ~Beach (Number of levels: 9)
##
```

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sd(Intercept)	3.83	1.20	2.05	6.73	1.00	994	1671
sd(NAP)	1.86	0.89	0.29	3.92	1.00	1006	1174
cor(Intercept,NAP)	-0.62	0.34	-0.99	0.28	1.00	1951	2401

```
##
## Population-Level Effects:
```

```
##           Estimate Est.Error 1-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept      6.63      1.36      3.91      9.30 1.00      1208      1671
## NAP           -2.71      0.83     -4.39     -1.07 1.00      1682      2019
##
## Family Specific Parameters:
##           Estimate Est.Error 1-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma         2.73      0.41      2.06      3.65 1.00      1226      2086
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

Again, instead of fitting all intercepts a_j and slopes b_j `brm()` fits the differences $\alpha_j = a_j - \mu_a$ and $\beta_j = b_j - \mu_b$. The model reads

$$y_i \sim \text{normal}((\mu_a + \alpha_{\text{group}(i)} + (\mu_b + \beta_{\text{group}(i)}) \cdot x, \sigma), \quad i = 1, \dots, n \quad (n \text{ observations})$$

$$\alpha_j \sim \text{normal}(0, \sigma_a), \quad j = 1, \dots, m \quad (m \text{ groups})$$


$$\beta_j \sim \text{normal}(0, \sigma_b), \quad j = 1, \dots, m \quad (m \text{ groups})$$

```
ranef(fit.b.part)
```

```
## $Beach
## , , Intercept
##
##      Estimate Est.Error      Q2.5      Q97.5
## 1  2.540847   1.743882 -0.8419498  6.0209580
## 2  5.788072   1.833344  2.3432109  9.3839597
## 3 -2.804893   1.683489 -6.0873369  0.4444730
## 4 -2.926661   1.800657 -6.4835468  0.5165117
## 5  4.207380   2.045448  0.1887037  8.2839585
## 6 -2.038257   1.654506 -5.3573790  1.2476575
## 7 -2.323106   1.935573 -6.3232722  1.4981540
## 8 -1.432882   1.745339 -4.7384368  2.0101851
## 9 -0.329528   1.759999 -3.7325519  3.1625538
##
## , , NAP
##
##      Estimate Est.Error      Q2.5      Q97.5
## 1  0.14891816  1.255007 -2.074871  3.01373291
## 2 -1.51070274  1.305144 -4.210327  0.86520672
## 3  0.88227617  1.140999 -1.300154  3.30303108
## 4  1.01779069  1.179102 -1.200473  3.52142295
## 5 -3.02565058  1.833075 -6.692674  0.02281932
## 6  0.92695921  1.071852 -1.093888  3.23927093
## 7  0.67046217  1.309826 -1.895710  3.41201102
## 8  0.58224544  1.046994 -1.325798  2.73604016
## 9 -0.07492481  1.098433 -2.286378  2.17586757
```

Now, let's look at observed and predicted. `conditional_effects()` plots fixed effects only by default. The whole dataset is shown and predictions with mean slope μ_b and mean intercept μ_a .

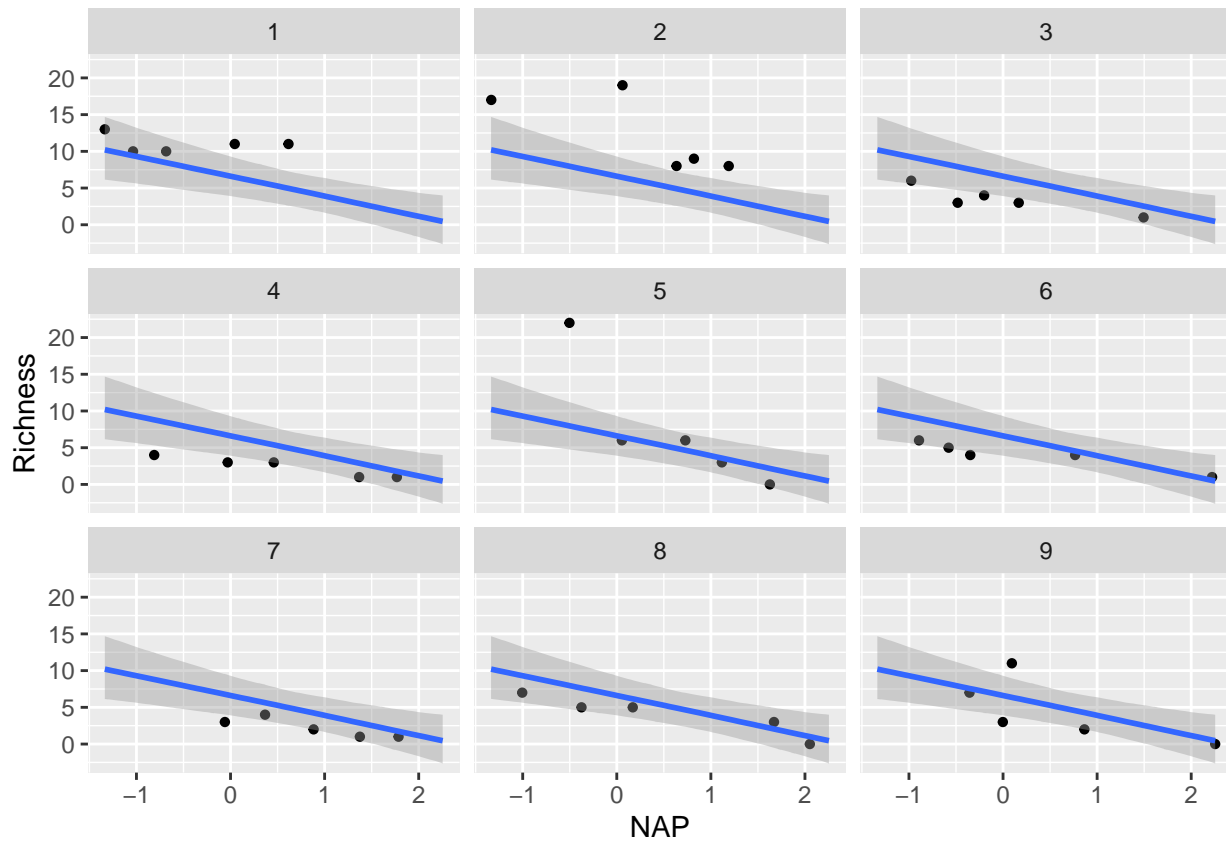
```
plot( conditional_effects(fit.b.part),
      points=TRUE,
      ask=FALSE )
```



When specifying group-level predictors (all levels) with `conditions=`, we receive a warning that they are not part of the (fixed effects) model. We see a plot for all levels of `Beach`, but the regression line is the same: predictions for fixed effects part only!

```
plot( conditional_effects(fit.b.part,
                          effects = "NAP",
                          conditions = data.frame(Beach=levels(df$Beach)) ),
      points=TRUE,
      ask=FALSE )
```


```
## Warning: The following variables in 'conditions' are not part of the model:
## 'Beach'
```



To include also the random effects for model predictions, `re_formula=NULL` must be specified. This reads a little weird (I would have expected something like `re_formula=TRUE`), but the default for no random effects as above is `re_formula=NA` and `NULL` is the command for using random effects for prediction here.

Slopes are identical, but intercepts vary between groups.

```
plot( conditional_effects(fit.b.part,
  effects = "NAP",
  re_formula = NULL,
  conditions = data.frame(Beach=levels(df$Beach)) ),
  points=TRUE,
  ask=FALSE )
```

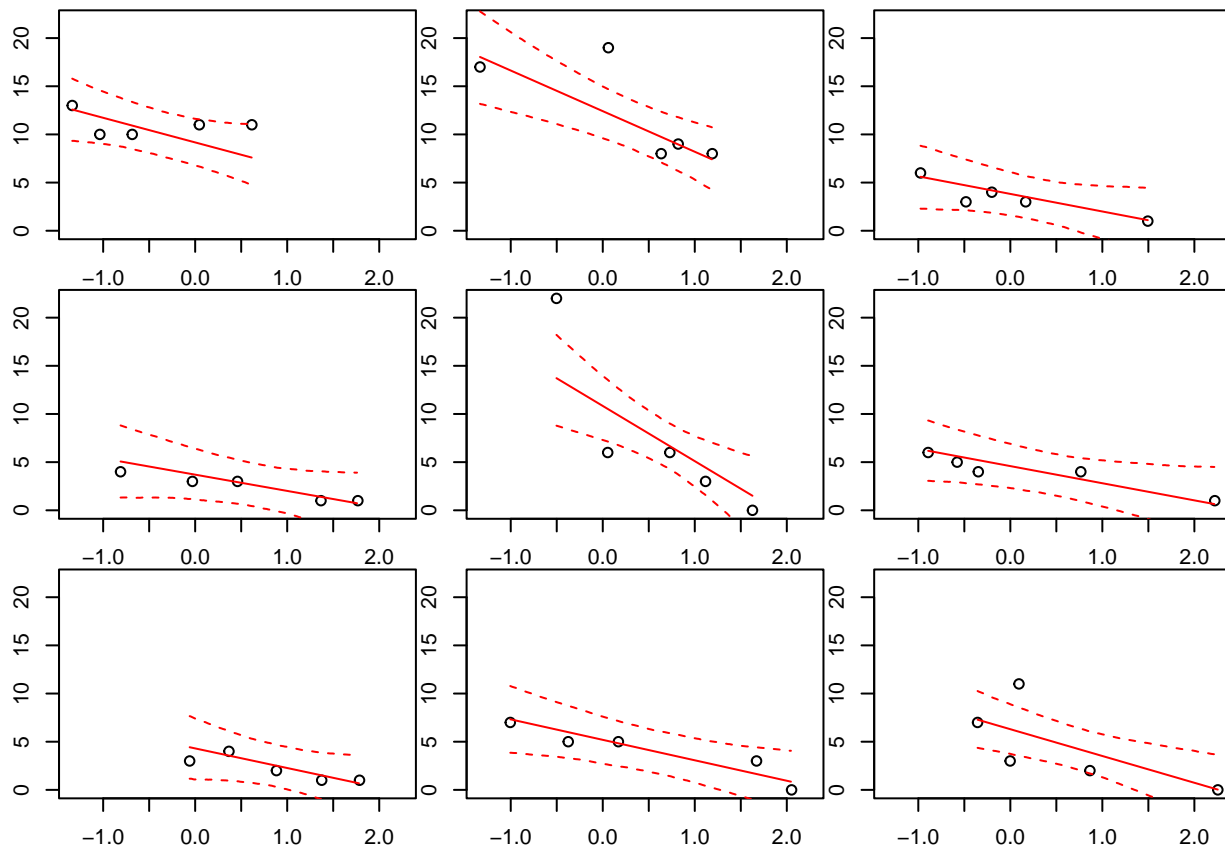
Same as in the “no pooling” model, group-level predictions can also be extracted by hand using `fitted()`.

```
levels = levels(df$Beach)
```

```
par(mfrow=c(3,3), mar=c(1,1,1,1), oma=c(1,1,0,0))
for (i in 1:9){
  df.sub=subset(df, df$Beach==i)
  x.pred = seq(from=min(df.sub$NAP), to=max(df.sub$NAP), by=0.01)
  plot(df.sub$NAP, df.sub$Richness,
       xlim=range(df$NAP),
       ylim=range(df$Richness) )

  y.cred = fitted(fit.b.part, newdata=data.frame(NAP=x.pred,
                                                Beach=levels[i] ) )

  lines(x.pred, y.cred[, 1], col="red")
  lines(x.pred, y.cred[, 3], col="red", lty=2)
  lines(x.pred, y.cred[, 4], col="red", lty=2)
}
```



We can plot the whole dataset with all group-level mean regression lines, both for the previous “no pooling” model and the current “partial pooling” model. With partial pooling, slopes b_j are drawn towards the overall mean intercept μ_b (i.e. less extreme) compared to no pooling.

```
plot(df$NAP, df$Richness)
```

```
# partial pooling
for (i in 1:9){
  df.sub=subset(df, df$Beach==i)
  x.pred = seq(from=min(df.sub$NAP), to=max(df.sub$NAP), by=0.01)
  y.cred = fitted(fit.b.part, newdata=data.frame(NAP=x.pred,
                                                  Beach=levels[i] ) )

  lines(x.pred, y.cred[, 1], col="red")
}

x.pred = seq(from=min(df$NAP), to=max(df$NAP), by=0.01)
y.cred = fitted(fit.b.part, newdata=data.frame(NAP=x.pred,
                                              Beach=NA) )

# partial pooling mean
lines(x.pred, y.cred[, 1], col="red", lwd=3, lty=3)

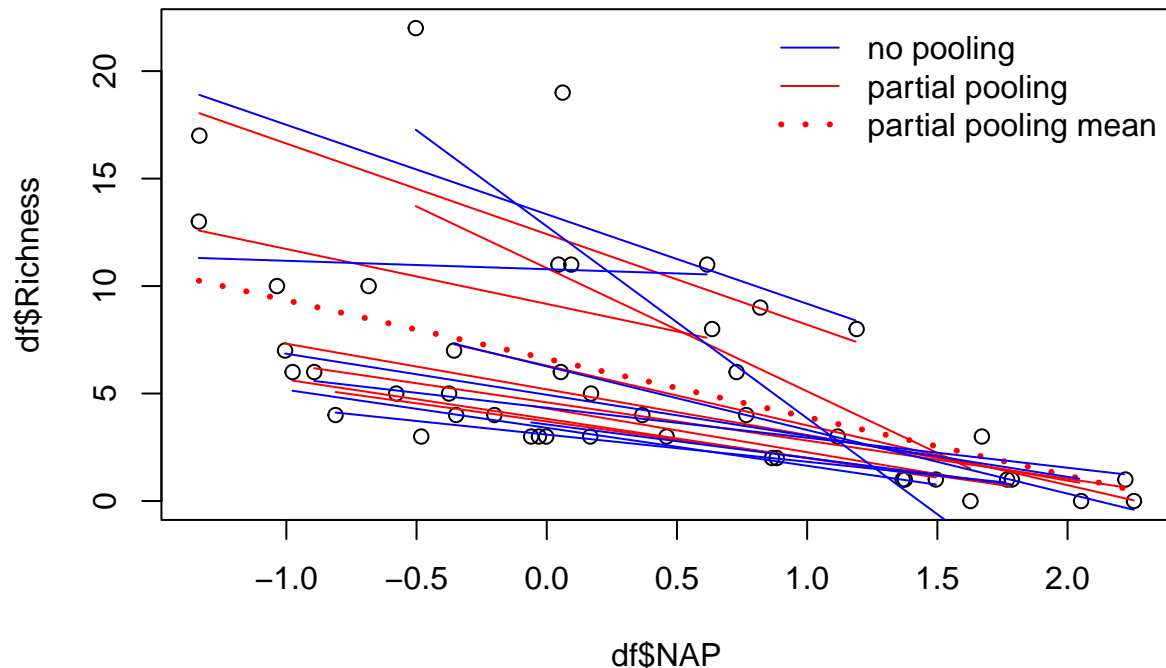
# no pooling (previous model)
for (i in 1:9){
  df.sub=subset(df, df$Beach==i)
  x.pred = seq(from=min(df.sub$NAP), to=max(df.sub$NAP), by=0.01)
  y.cred = fitted(fit.b.no, newdata=data.frame(NAP=x.pred,
```

```

                                Beach=levels[i] ) )
  lines(x.pred, y.cred[, 1], col="blue")
}

legend("topright",
      legend=c("no pooling","partial pooling", "partial pooling mean"),
      lwd=c(1,1,3),
      lty=c(1,1,3),
      col=c("blue","red","red"),
      bty="n")

```



Extra: GLMM

We assumed normally distributed residuals and the default in `brms` is Gaussian indeed. But the response we model (species richness) is an integer (not continuous), it's also non-negative. So we could use a Poisson distribution for the stochastic part of the model, which is standard practice for count data. By specifying `family=poisson`, also a log-link is chosen by default. This means the regression lines become an exponential functions, which are non-negative, too.

```

fit.b.GLMM = brm( Richness ~ 1+NAP + (1+NAP|Beach),
                  family=poisson,
                  data=df )

```

```
fit.b.GLMM
```

```


## Warning: There were 16 divergent transitions after warmup. Increasing
## adapt_delta above 0.8 may help. See http://mc-stan.org/misc/
## warnings.html#divergent-transitions-after-warmup

## Family: poisson
## Links: mu = log
## Formula: Richness ~ 1 + NAP + (1 + NAP | Beach)
## Data: df (Number of observations: 45)

```

```
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
## total post-warmup draws = 4000
##
## Group-Level Effects:
## ~Beach (Number of levels: 9)
##      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sd(Intercept)      0.70     0.24    0.38    1.29 1.00    1315    2129
## sd(NAP)            0.45     0.22    0.14    0.98 1.00     908     906
## cor(Intercept,NAP) 0.11     0.39   -0.67    0.81 1.00    2094    2171
##
## Population-Level Effects:
##      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept      1.68     0.26    1.16    2.18 1.00    1203    1597
## NAP           -0.62     0.20   -1.03   -0.24 1.01    1236     851
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

```
plot( conditional_effects(fit.b.GLMM),
      points=TRUE,
      ask=FALSE )
```



```
p = plot( conditional_effects(fit.b.GLMM,
                             effects = "NAP",
                             re_formula = NULL,
                             conditions = data.frame(Beach=levels(df$Beach)) ),
          points=TRUE,
```

```
ask=FALSE,
plot=FALSE)
p[[1]] + ggplot2::ylim(c(0,25))
```

