

Introduction to Bayesian Statistics

Part 6 Linear Mixed Models

Benjamin Rosenbaum

iDiv 2025



This lecture

- Random grouping factors
- Random effects
- Fixed & random effects
 - Random intercepts
 - Random intercepts and slopes
- Multiple random grouping factors
- Continuous correlation structures
- Summary

Assumptions for linear models

1. Independent observations.

Systematic differences in y are because of x !

2. Trend of y follows (linear) prediction model


$$\mu(x) = a + b \cdot x$$

3. Residuals follow normal distribution

$$\varepsilon \sim \text{Normal}(0, \sigma)$$

4. Constant variance (standard deviation σ)

across whole range of x




Assumptions for linear models

When are residuals not independent ?

Study design: N observations have been **sampled in groups**


- There is a categorical predictor with M levels ($M < N$)
- Each level / group contains multiple observations, these are no independent replicates !
- Maybe structural (or random) differences between (unobserved) conditions of these groups
- Still want to make **inference on whole population**



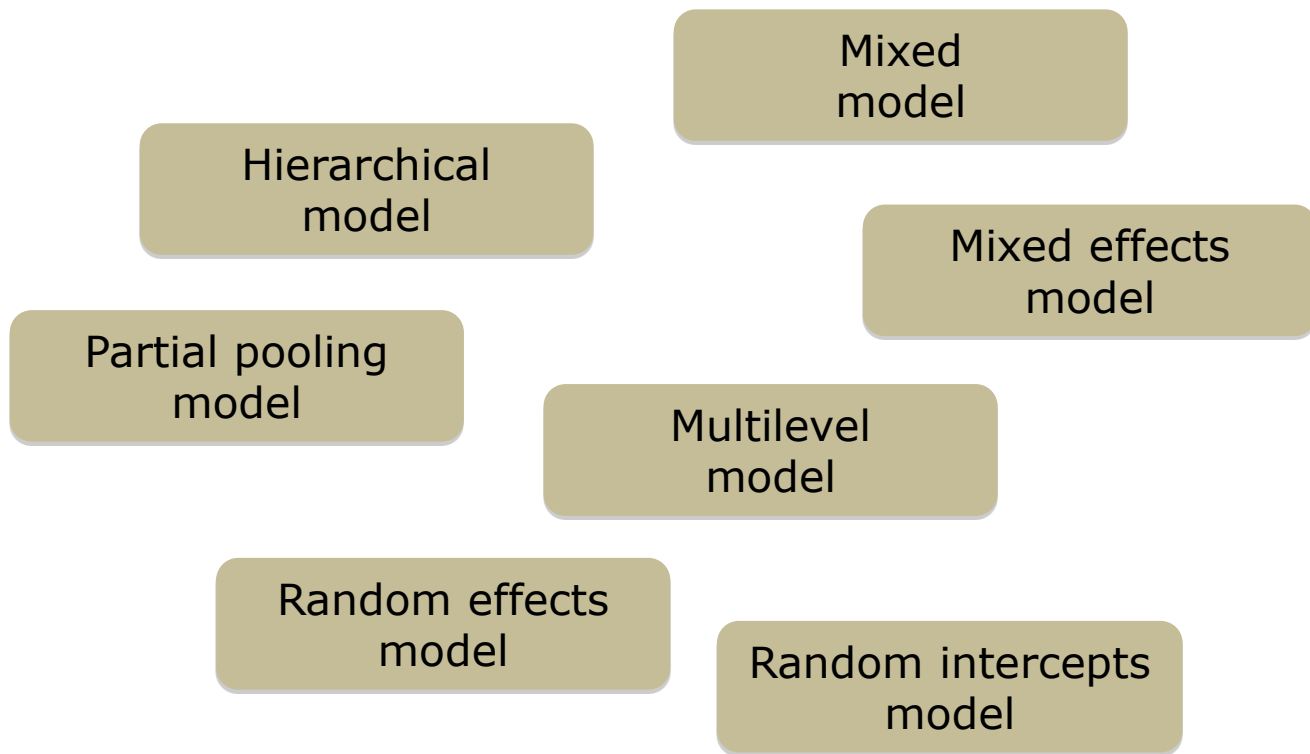
Random grouping factors

Examples:

- Site: Data collected at multiple locations
- Year: Data collected ~same time in different years
- Observer: Data collected by different researchers
- Individual: Multiple measurements on the same test object
- Unit: Multiple experimental units like chambers / tanks
- Species: Data on higher level (e.g. family) contains multiple species



Word salad !!



Random effects

Example

- Measured weight of $N = 200$ viper snakes
- Data collected on $M = 9$ different sites across France
- Also measured length, but we're not using it for now


Q: What is the correct statistical model for species mean weight ?

`brm(weight ~ 1) ?`

→ No, observations not independent

`brm(weight ~ site) ?`

→ No, does not estimate overall weight (just site-specific means)




(1) Complete pooling

Model: $y_i = a + \varepsilon_i$ $i = 1 \dots N$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

Priors: $a \sim \text{Normal}(150, 10)$

- Ignore categorical predictor *site*
- Information is completely pooled across all levels
- Fit joint intercept a only
- a is given a prior




(2) No pooling

Model: $y_i = a_{site(i)} + \varepsilon_i \quad i = 1 \dots N$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

Priors: $a_j \sim \text{Normal}(150, 10) \quad j = 1 \dots M$

- Include categorical predictor *site* (**ANOVA**)
- No information is shared across levels
- Fit independent intercepts $a_1 \dots a_M$
- Each a_j is given a prior
(here they are identical, but could also be different)



(3) Partial pooling

Model: $y_i = a_{\text{site}(i)} + \varepsilon_i \quad i = 1 \dots N$

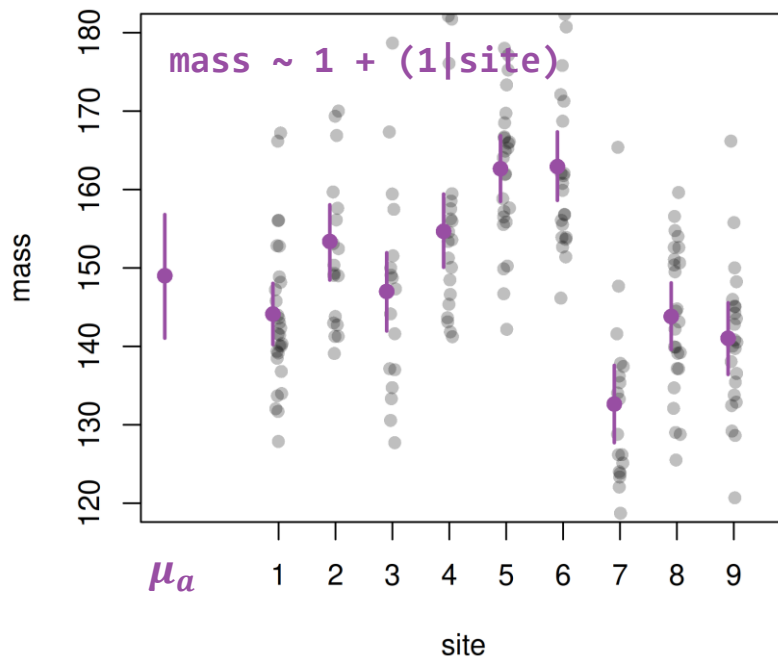
$\varepsilon_i \sim \text{Normal}(0, \sigma)$

$a_j \sim \text{Normal}(\mu_a, \sigma_a) \quad j = 1 \dots M$

Priors: $\mu_a \sim \text{Normal}(150, 10)$

$\sigma_a \sim \text{Exponential}(1/10)$

- Replace prior distributions by a joint distribution of the site-specific intercepts $a_1 \dots a_M$
- Their mean μ_a and sdev σ_a are parameters and are estimated with all others by MCMC
- These “hyperparameters” μ_a, σ_a are given priors



(3) Partial pooling

Model: $y_i = \mu_a + \delta_{\text{site}(i)} + \varepsilon_i \quad i = 1 \dots N$


$\varepsilon_i \sim \text{Normal}(0, \sigma)$

$\delta_j \sim \text{Normal}(0, \sigma_a) \quad j = 1 \dots M$

Priors: $\mu_a \sim \text{Normal}(150, 10)$

$\sigma_a \sim \text{Exponential}(1/10)$

- Replaced $a_j = \mu_a + \delta_j$
- Effects $\delta_j = a_j - \mu_a$ describe each site's deviation from overall mean μ_a
- Identical model, just rearrangement of coefficients



Comparison

vs. complete pooling $\text{mass} \sim 1$

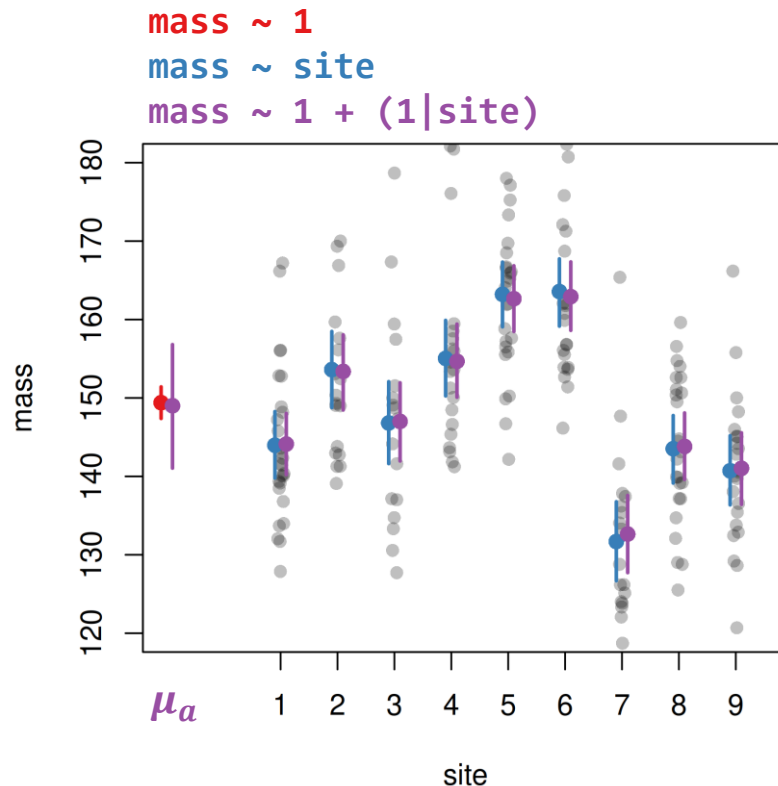
Uncertainty of μ_a in $\text{mass} \sim 1$ is smaller.
Wrongly assumes $N = 200$ independent obs.

Uncertainty of μ_a in $\text{mass} \sim 1 + (1|\text{site})$ correct.
Informed by $M = 9$ sites.

vs. no pooling $\text{mass} \sim \text{site}$

Site-specific means a_j of $\text{mass} \sim 1 + (1|\text{site})$
closer to overall mean μ_a than in $\text{mass} \sim \text{site}$

„Shrinkage“ of parameters



Random effects


Random effects $\delta_j = a_j - \mu_a$ describe each site's deviation from overall mean μ_a with sdev σ_a

Assumption: normally distributed

(Bayesian 3d printer: could also be other distribution, e.g. positive lognormal distribution for a_j)

Hyperparameters μ_a, σ_a informed by all 9 sites

- Some information is shared / pooled
- Borrowing strength
- For unbalanced designs this improves prediction in levels with few observations.



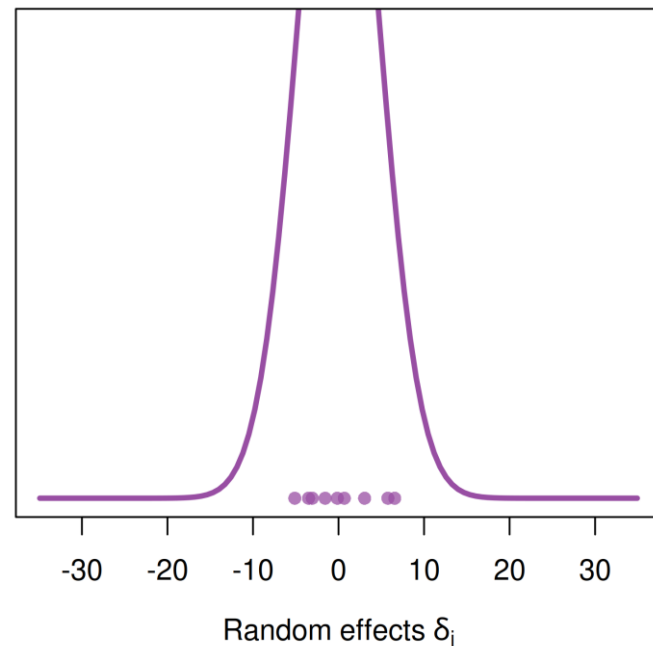
Random effects

Random effects $\delta_j = a_j - \mu_a$ describe each site's deviation from overall mean μ_a with sdev σ_a

This part of the likelihood $\delta_j \sim \text{Normal}(0, \sigma_a)$ would be maximized for $\sigma_a \rightarrow 0$ and all $\delta_j \rightarrow 0$

But the residuals' likelihood $y_i \sim \text{Normal}(a_{\text{site}(i)}, \sigma)$ would decrease if all δ_j were close to 0

→ Partial pooling model is a compromise between small random effects (on site-level) and good model fit (on observation level)



Variance partitioning

Overall variance of the data is estimated on different levels


$$s_{\text{total}}^2 = s_{\text{among}}^2 + s_{\text{within}}^2$$

Site-level (among) σ_a

Explains site-specific deviation from overall mean μ_a

Observation-level (within) σ

Explains each data point's deviation from site-mean a_j
(All the randomness in the data which is not caused by site-level differences)



> brm(mass ~ 1 + (1|site), ...)

```
Family: gaussian
Links: mu = identity; sigma = identity
Formula: mass ~ 1 + (1 | site)
Data: df (Number of observations: 200)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000
```

Multilevel Hyperparameters:

~site (Number of levels: 9)

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
$\sigma_a \rightarrow$ sd(Intercept)	11.64	3.26	6.99	19.70	1.01	814	1353

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
$\mu_a \rightarrow$ Intercept	149.01	3.91	141.05	156.81	1.00	777	1038

Further Distributional Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
$\sigma \rightarrow$ sigma	11.02	0.57	9.97	12.20	1.00	2527	2566

Fixed and random effects

Fixed & random effects


In previous example, we were interested in a correct statistical model for the overall intercept

→ „**Fixed effect**“ μ_a

Variation among sites was not our focus, but we wanted to account for non-independence of observations

→ „**Random effect**“ δ_j ($j = 1 \dots M$)

Fitted an intercept-only model with random intercepts (site)



Example

- Measured weight of $N = 200$ viper snakes
- Data collected on $M = 9$ different sites across France
- **Also measured length (continuous predictor)**


Q: Correct statistical regression model for weight vs. length?

`brm(weight ~ length) ?`

→ No, observations not independent

`brm(weight ~ length*site) ?`

→ No, does not estimate overall slope (just site-specific regressions)



(1) Complete pooling


Model: $y_i = a + b \cdot \text{length}_i + \varepsilon_i \quad i = 1 \dots N$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

Priors: $a \sim \text{Normal}(150, 10)$

$b \sim \text{Normal}(10, 5)$

- Ignore categorical predictor *site*
- Information is completely pooled across all levels
- Fit joint intercept a & slope b only
- a, b are given a prior



(2) No pooling


Model: $y_i = a_{site(i)} + b_{site(i)} \cdot length_i + \varepsilon_i$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

Priors: $a_j \sim \text{Normal}(150, 10)$ $j = 1 \dots M$

$b_j \sim \text{Normal}(10, 5)$

- Include categorical predictor *site* (**ANCOVA**)
- No information is shared across levels
- Fit independent intercepts $a_1 \dots a_M$ & slopes $b_1 \dots b_M$
- Each a_j & b_j is given a prior
(here they are identical, but could also be different)



(2) No pooling

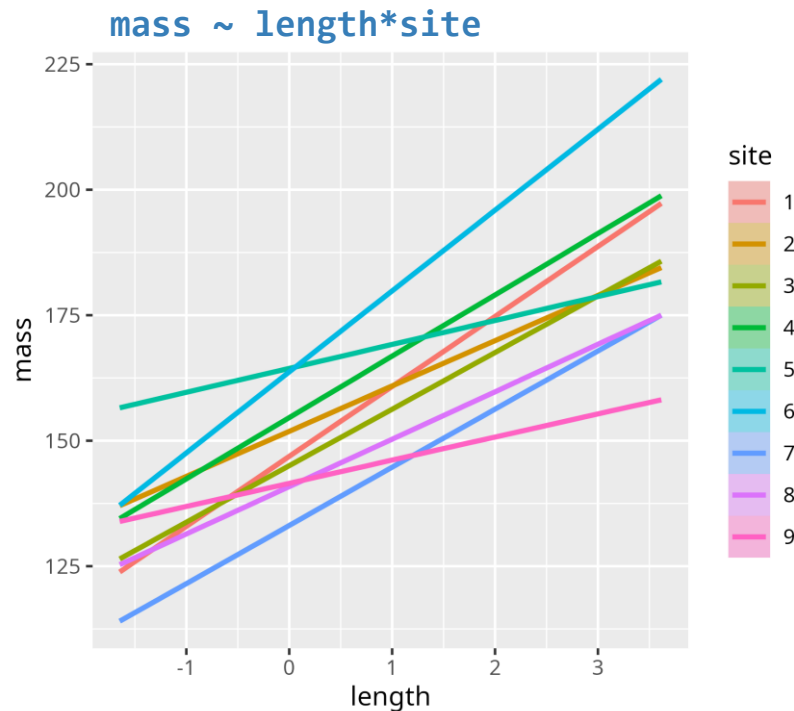
Model: $y_i = a_{site(i)} + b_{site(i)} \cdot length_i + \varepsilon_i$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

Priors: $a_j \sim \text{Normal}(150, 10)$ $j = 1 \dots M$

$b_j \sim \text{Normal}(10, 5)$

- Include categorical predictor *site* (**ANCOVA**)
- No information is shared across levels
- Fit independent intercepts $a_1 \dots a_M$ & slopes $b_1 \dots b_M$
- Each a_j & b_j is given a prior
(here they are identical, but could also be different)



(3) Partial pooling: Random intercepts

Model: $y_i = \mathbf{a}_{\text{site}(i)} + b \cdot \text{length}_i + \varepsilon_i$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

$\mathbf{a}_j \sim \text{Normal}(\mu_a, \sigma_a) \quad j = 1 \dots M$

Priors: $b_j \sim \text{Normal}(10, 5)$


$\mu_a \sim \text{Normal}(150, 10)$

$\sigma_a \sim \text{Exponential}(1/10)$

Random effects: $a_1 \dots a_M$ site-specific intercepts

Fixed effects: μ_a joint intercept

b joint slope



Marginal predictions: $\mu_a + b \cdot \text{length}$

(3) Partial pooling: Random intercepts

Model: $y_i = (\mu_a + \delta_{\text{site}(i)}) + b \cdot \text{length}_i + \varepsilon_i$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

$\delta_j \sim \text{Normal}(0, \sigma_a) \quad j = 1 \dots M$

Priors: $b_j \sim \text{Normal}(10, 5)$


$\mu_a \sim \text{Normal}(150, 10)$

$\sigma_a \sim \text{Exponential}(1/10)$

Random effects: $\delta_1 \dots \delta_M$ intercepts deviation

Fixed effects: μ_a joint intercept

b joint slope



Marginal predictions: $\mu_a + b \cdot \text{length}$

(3) Partial pooling: Random intercepts

Model: $y_i = (\mu_a + \delta_{\text{site}(i)}) + b \cdot \text{length}_i + \varepsilon_i$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

$\delta_j \sim \text{Normal}(0, \sigma_a)$ $j = 1 \dots M$

Priors: $b_j \sim \text{Normal}(10, 5)$

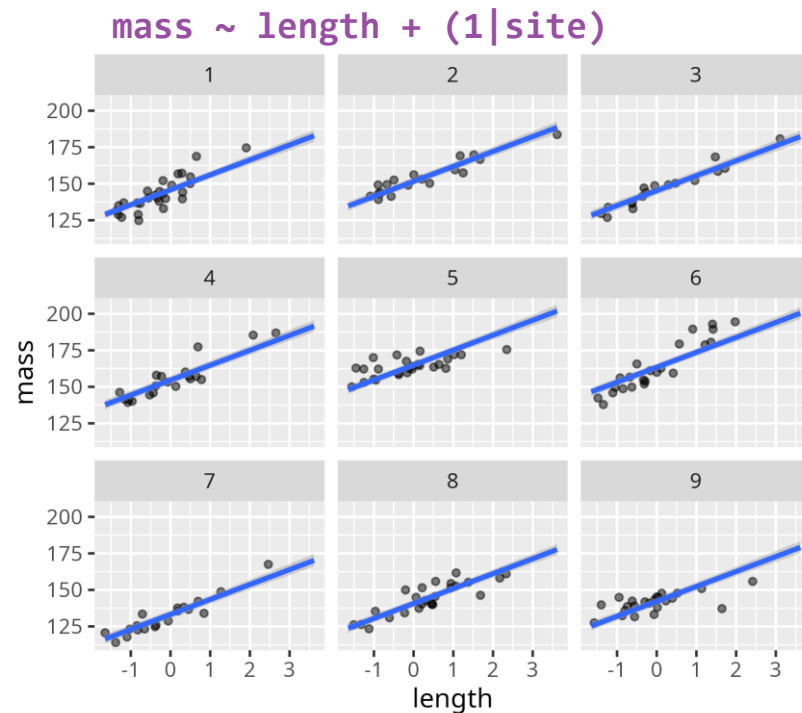
$\mu_a \sim \text{Normal}(150, 10)$

$\sigma_a \sim \text{Exponential}(1/10)$

Random effects: $\delta_1 \dots \delta_M$ intercepts deviation

Fixed effects: μ_a joint intercept

b joint slope



Conditional predictions: $a_j + b \cdot \text{length}$

(3) Partial pooling: Random intercepts

Model: $y_i = (\mu_a + \delta_{\text{site}(i)}) + b \cdot \text{length}_i + \varepsilon_i$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

$\delta_j \sim \text{Normal}(0, \sigma_a) \quad j = 1 \dots M$

Priors: $b_j \sim \text{Normal}(10, 5)$

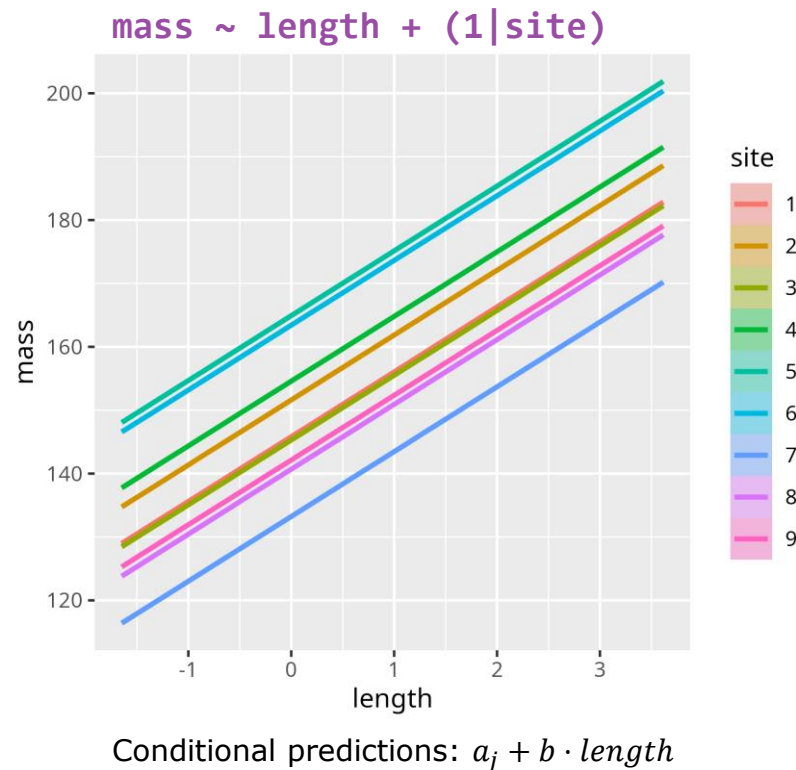
$\mu_a \sim \text{Normal}(150, 10)$

$\sigma_a \sim \text{Exponential}(1/10)$

Random effects: $\delta_1 \dots \delta_M$ intercepts deviation

Fixed effects: μ_a joint intercept

b joint slope



(4) Partial pooling: Random slopes & intercepts

Model: $y_i = a_{\text{site}(i)} + b_{\text{site}(i)} \cdot \text{length}_i + \varepsilon_i$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

$a_j \sim \text{Normal}(\mu_a, \sigma_a) \quad j = 1 \dots M$

$b_j \sim \text{Normal}(\mu_b, \sigma_b)$

Priors $\mu_a \sim \text{Normal}(150, 10)$


$\mu_b \sim \text{Normal}(10, 5)$

$\sigma_a \sim \text{Exponential}(1/10)$

$\sigma_b \sim \text{Exponential}(1/5)$

Random effects: $a_1 \dots a_M, b_1 \dots b_M$

Fixed effects: μ_a, μ_b joint intercept & slope



Marginal predictions: $\mu_a + \mu_b \cdot \text{length}$

(4) Partial pooling: Random slopes & intercepts

Model: $y_i = (\mu_a + \delta_{\text{site}(i)}) + (\mu_b + \gamma_{\text{site}(i)}) \cdot \text{length}_i + \varepsilon_i$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

$\delta_j \sim \text{Normal}(0, \sigma_a) \quad j = 1 \dots M$

$\gamma_j \sim \text{Normal}(0, \sigma_b)$

Priors $\mu_a \sim \text{Normal}(150, 10)$

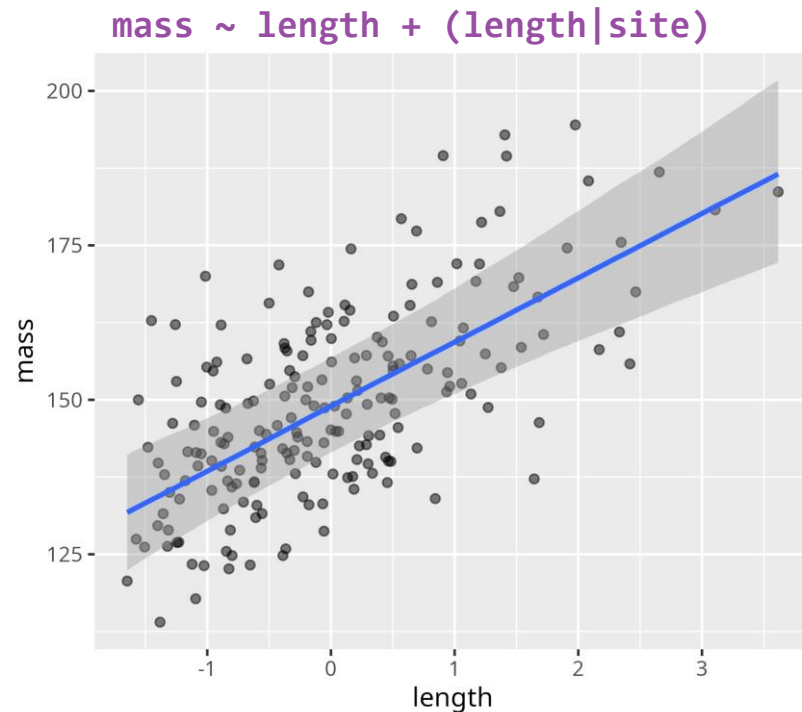
$\mu_b \sim \text{Normal}(10, 5)$

$\sigma_a \sim \text{Exponential}(1/10)$

$\sigma_b \sim \text{Exponential}(1/5)$

Random effects: $\delta_1 \dots \delta_M, \gamma_1 \dots \gamma_M$

Fixed effects: μ_a, μ_b joint intercept & slope



Marginal predictions: $\mu_a + \mu_b \cdot \text{length}$

(4) Partial pooling: Random slopes & intercepts

Model: $y_i = (\mu_a + \delta_{site(i)}) + (\mu_b + \gamma_{site(i)}) \cdot length_i + \varepsilon_i$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

$\delta_j \sim \text{Normal}(0, \sigma_a) \quad j = 1 \dots M$

$\gamma_j \sim \text{Normal}(0, \sigma_b)$

Priors $\mu_a \sim \text{Normal}(150, 10)$


$\mu_b \sim \text{Normal}(10, 5)$

$\sigma_a \sim \text{Exponential}(1/10)$

$\sigma_b \sim \text{Exponential}(1/5)$

Random effects: $\delta_1 \dots \delta_M, \gamma_1 \dots \gamma_M$

Fixed effects: μ_a, μ_b joint intercept & slope



Conditional predictions: $a_j + b_j \cdot length$

(4) Partial pooling: Random slopes & intercepts

Model: $y_i = (\mu_a + \delta_{\text{site}(i)}) + (\mu_b + \gamma_{\text{site}(i)}) \cdot \text{length}_i + \varepsilon_i$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

$\delta_j \sim \text{Normal}(0, \sigma_a) \quad j = 1 \dots M$

$\gamma_j \sim \text{Normal}(0, \sigma_b)$

Priors $\mu_a \sim \text{Normal}(150, 10)$


$\mu_b \sim \text{Normal}(10, 5)$

$\sigma_a \sim \text{Exponential}(1/10)$

$\sigma_b \sim \text{Exponential}(1/5)$

Random effects: $\delta_1 \dots \delta_M, \gamma_1 \dots \gamma_M$

Fixed effects: μ_a, μ_b joint intercept & slope



Conditional predictions: $a_j + b_j \cdot \text{length}$

(4) Partial pooling: Random slopes & intercepts

Model: $y_i = (\mu_a + \delta_{site(i)}) + (\mu_b + \gamma_{site(i)}) \cdot length_i + \varepsilon_i$
 $\varepsilon_i \sim \text{Normal}(0, \sigma)$

$$\begin{pmatrix} \delta_j \\ \gamma_j \end{pmatrix} \sim \text{MVNormal} \left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \begin{bmatrix} \sigma_a^2 & \rho \sigma_a \sigma_b \\ \rho \sigma_a \sigma_b & \sigma_b^2 \end{bmatrix} \right)$$

$j = 1 \dots M$


Priors $\mu_a \sim \text{Normal}(150, 10)$
 $\mu_b \sim \text{Normal}(10, 5)$
 $\sigma_a \sim \text{Exponential}(1/10)$
 $\sigma_b \sim \text{Exponential}(1/5)$
 $\rho \sim \text{Uniform}(0, 1)$

← Random intercepts & slopes not independent:


Multivariate normal distributions with zero mean and 2x2 covariance matrix Σ

← Correlation coefficient ρ

Comparison



Uncertainty of a, b is smaller (overly confident).
Wrongly assumes $N = 200$ independent obs.

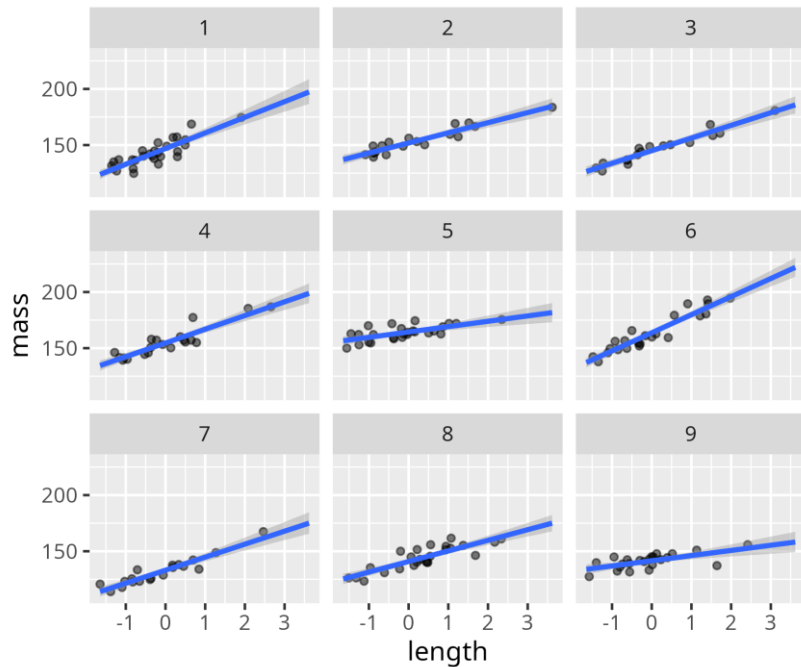


Uncertainty of μ_a, μ_b is correct.

Comparison

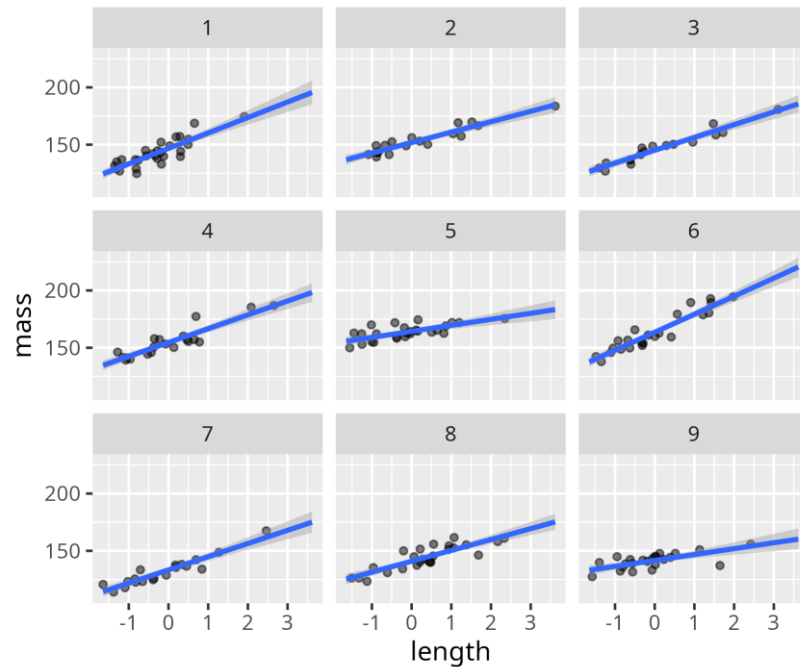
Here,
difference is
very weak

$\text{mass} \sim \text{length} * \text{site}$



Regressions for each site independent.

$\text{mass} \sim \text{length} + (\text{length} | \text{site})$




Intercepts & slopes draws to joint means
„Shrinkage“

Comparison


Here,
difference is
very weak

$\text{mass} \sim \text{length} * \text{site}$




Regressions for each site independent.

$\text{mass} \sim \text{length} + (\text{length} | \text{site})$




Intercepts & slopes draws to joint means
„Shrinkage“

Comparison: unbalanced design



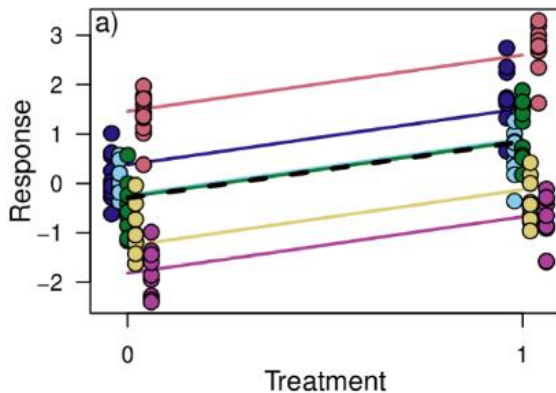
High uncertainty in sites with few data.




Borrowing strength from other sites.
→ Better predictions!

**Categorical
predictor**



Random intercepts



Random intercepts & „slopes“



**Continuous
predictor**



Harrison et al
(2020) *PeerJ*

```
> brm(mass ~ length + (length|site), ...)
```

```
Family: gaussian
Links: mu = identity; sigma = identity
Formula: mass ~ length + (length | site)
Data: df (Number of observations: 200)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000
```

Multilevel Hyperparameters:

~site (Number of levels: 9)

		Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS
$\sigma_a \rightarrow$	sd(Intercept)	12.36	3.45	7.39	20.97	1.00	1269
$\sigma_b \rightarrow$	sd(length)	4.65	1.57	2.52	8.72	1.00	1154
$\rho \rightarrow$	cor(Intercept,length)	0.05	0.33	-0.58	0.64	1.00	1937

Regression Coefficients:

		Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
$\mu_a \rightarrow$	Intercept	148.95	3.88	141.48	156.81	1.00	1265	2055
$\mu_b \rightarrow$	length	10.43	1.69	7.11	13.81	1.00	1318	1461

Further Distributional Parameters:

		Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
$\sigma \rightarrow$	sigma	5.28	0.27	4.79	5.86	1.00	3411	2993


Multiple grouping factors

Example

- Measured weight of $N = 200$ viper snakes
- Data collected on $M = 9$ different sites across France
- **Collection was done in $L = 5$ sampling campaigns**

Can be fully crossed (each campaign visited all sites)
or partially crossed (not all campaigns visited all sites).

→ Specify **additive (independent)** random effects
for grouping factor site and grouping factor campaign



Crossed random effects

Model: $y_i = \mu_a + \delta_{site(i)} + \gamma_{campaign(i)} + \varepsilon_i$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

$\delta_j \sim \text{Normal}(0, \sigma_\delta) \quad j = 1 \dots M$

$\gamma_k \sim \text{Normal}(0, \sigma_\gamma) \quad k = 1 \dots L$

Priors: $\mu_a \sim \text{Normal}(150, 10)$

$\sigma_\delta \sim \text{Exponential}(1/10)$

$\sigma_\gamma \sim \text{Exponential}(1/10)$

$\text{mass} \sim 1 + (1|\text{site}) + (1|\text{campaign})$

← Site-effect

← Campaign-effect

← Overall mean is the only fixed effect in this basic example

Variance partitioning

Overall variance of the data is estimated on different levels

$$s_{\text{total}}^2 = s_{\text{among}}^2 + s_{\text{within}}^2$$


Site-campaign-level (among) $\sigma_{\delta} + \sigma_{\gamma}$

Explains additive deviations from overall mean μ_a
caused by site (σ_{δ}) and campaign (σ_{γ})

Observation-level (within) σ

Explains each data point's deviation from
site-campaign-mean $\mu_a + \delta_j + \gamma_k$

(All the randomness in the data which is not caused by site-level differences)



> brm(mass ~ 1 + (1|site) + (1|campaign), ...)

Family: gaussian

Links: mu = identity; sigma = identity

Formula: mass ~ 1 + (1 | site) + (1 | campaign)

Data: df (Number of observations: 200)

Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;

total post-warmup draws = 4000

Multilevel Hyperparameters:

~campaign (Number of levels: 5)

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
$\sigma_\gamma \rightarrow$ sd(Intercept)	1.39	1.33	0.05	4.94	1.00	1983	2233

~site (Number of levels: 9)

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
$\sigma_\delta \rightarrow$ sd(Intercept)	11.38	3.22	6.89	19.27	1.01	1085	2004

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
$\mu_a \rightarrow$ Intercept	149.14	3.86	141.37	156.68	1.00	828	1274

Further Distributional Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
$\sigma \rightarrow$ sigma	11.04	0.58	9.99	12.20	1.00	3320	2678


Example

- Measured weight of $N = 200$ viper snakes
- Data collected on $M = 9$ different sites across France
- **Sites belong to $L = 3$ distinct mountain regions**

Every site uniquely belongs to one region.

At least one region has 2 or more sites.

→ Specify **nested** random effects
for factor site in factor region



Nested random effects

Model: $y_i = \mu_a + \gamma_{region(i)} + \delta_{region:site(i)} + \varepsilon_i$

$\varepsilon_i \sim \text{Normal}(0, \sigma)$

$\gamma_k \sim \text{Normal}(0, \sigma_\gamma) \quad k = 1 \dots L$

$\delta_j \sim \text{Normal}(0, \sigma_\delta) \quad j = 1 \dots M$

Priors: $\mu_a \sim \text{Normal}(150, 10)$

$\sigma_\gamma \sim \text{Exponential}(1/10)$

$\sigma_\delta \sim \text{Exponential}(1/10)$

$\text{mass} \sim 1 + (1|\text{region/site})$

$\text{mass} \sim 1 + (1|\text{region}) + (1|\text{region:site})$

← Regional-effect

← Site-effect in region

Automatically
uses unique
region:site
combinations
with M levels

Variance partitioning

Overall variance of the data is estimated on different levels

$$s_{\text{total}}^2 = s_{\text{among regions}}^2 + s_{\text{among sites}}^2 + s_{\text{within}}^2$$

Regional-level σ_γ

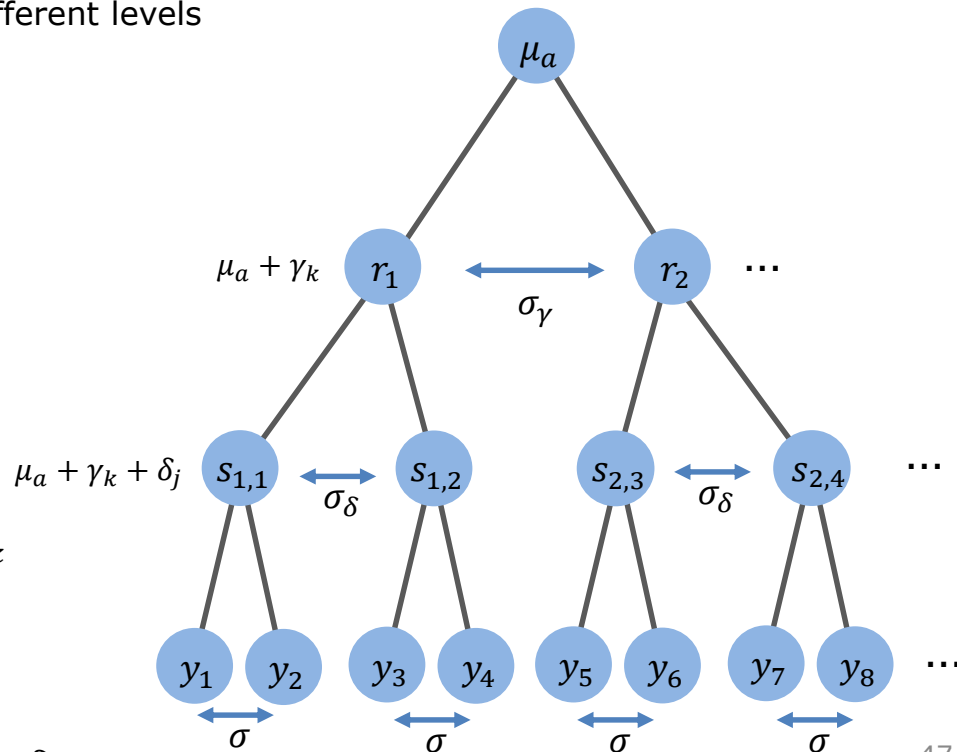
Explains deviations from overall mean μ_a

Site-level σ_δ

Explains deviations from regional mean $\mu_a + \gamma_k$

Observation-level (within) σ

Data point deviations from site-mean $\mu_a + \gamma_k + \delta_j$



```
> brm(mass ~ 1 + (1|region/site), ...)
> brm(mass ~ 1 + (1|region) + (1|region:site), ...)
```

```
Family: gaussian
Links: mu = identity; sigma = identity
Formula: mass ~ 1 + (1 | region/site)
Data: df (Number of observations: 200)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000
```

Multilevel Hyperparameters:

~region (Number of levels: 3)

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
$\sigma_\gamma \rightarrow$ sd(Intercept)	12.21	6.76	2.81	30.38	1.00	1155	1073

~region:site (Number of levels: 9)

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
$\sigma_\delta \rightarrow$ sd(Intercept)	6.01	2.60	2.57	12.65	1.00	1003	1395

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
$\mu_a \rightarrow$ Intercept	148.99	6.45	135.11	161.86	1.00	1594	1365

Further Distributional Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
$\sigma \rightarrow$ sigma	11.05	0.57	10.01	12.21	1.00	3750	2543

Nested or crossed?

Nested designs

$(1|F1/F2)$ just short for $(1|F1)+(1|F1:F2)$

If nested factor **F2** is uniquely labelled

(levels of **F2** and **F1:F2** identical)

this is same as crossed $(1|F1)+(1|F2)$

Crossed designs

No need to be fully crossed for additive model.

$(1|F1)+(1|F2)$ works for partially crossed, too.

Table 1. Schematic illustration of crossed and nested designs

Nested design $(1 F1/F2)$				
	Factor 2			
Factor 1	a	b	c	d
A	X	X		
B			X	X

Partially crossed design $(1 F1)+(1 F2)$				
	Factor 2			
Factor 1	a	b	c	d
A	X	X		
B	X	X		
C			X	X
D			X	X

Fully crossed design $(1 F1)+(1 F2)$				
	Factor 2			
Factor 1	a	b	c	d
A	X	X	X	X
B	X	X	X	X
C	X	X	X	X
D	X	X	X	X

Continuous correlation structures


Example

- Measured weight of $N = 200$ viper snakes
- Data collected on $M = 9$ different sites in the same region

Instead of using site as categorical grouping predictor, we could use location (longitude & latitude) as continuous „grouping“ predictor.

„Everything is related to everything else,
but near things are more related than distant things.“ (Tobler)

→ Specify spatial autocorrelation model.





Example

- Measured weight of $N = 200$ viper snakes
- Data collected on $M = 9$ different sites in the same region

Instead of using site as categorical grouping predictor, we could use location (longitude & latitude) as continuous „grouping“ predictor.

„Everything is related to everything else,
but near things are more related than distant things.“ (Tobler)

→ Specify spatial autocorrelation model.



Spatial autocorrelation model

Model: $y_i = \mu_a + \boldsymbol{\varepsilon}_i$

$$\boldsymbol{\varepsilon} \sim \text{MVNormal}(\mathbf{0}, \sigma^2 \mathbf{R})$$

Correlation: $R_{ij} = \exp(-\text{dist}(i, j)/\theta)$

Priors: $\mu_a \sim \text{Normal}(150, 10)$

$$\sigma \sim \text{Exponential}(1/10)$$

$$\theta \sim \text{Exponential}(1/100)$$

N -dimensional vector of residuals $\boldsymbol{\varepsilon}$:
multivariate normal distribution

← with $N \times N$ correlation matrix \mathbf{R} & variance σ^2

← Correlation decreases with
spatial distance of datapoints i, j

← Range parameter of spatial autocorrelation

- Model for spatial autocorrelation of **residuals**.
- Unfortunately not (yet?) implemented in brms.
- Code in **Stan**, or use alternative models.

Fixed residual correlation

Model: $y_i = \mu_a + \varepsilon_i$
 $\varepsilon \sim \text{MVNormal}(0, \sigma^2 R)$

Correlation: R_{ij} provided by user

Priors: $\mu_a \sim \text{Normal}(150, 10)$
 $\sigma \sim \text{Exponential}(1/10)$

```
brm( y ~ 1 + fcor(R),  
      data = ... ,  
      data2 = list(R=R) )
```

← Choose range θ yourself and compute, e.g.,
 $R_{ij} = \exp(-\text{dist}(i, j)/\theta)$ from data.
Obtain θ from variogram.

- Not ideal, since correlation range θ is not estimated jointly with the other parameters.
- But it's a fast and efficient solution.


More „adventures in covariance“ (McElreath)

brms has a couple of options:

- Fixed correlation model
- Gaussian process regression
- Generalized additive models (GAM)
→ but see mgcv package
- Adjacency matrix $R_{ij} = \begin{cases} 1 & i, j \text{ neighbours} \\ 0 & \text{otherwise} \end{cases}$
- ...

You can model autocorrelations

- spatial
- temporal
- spatio-temporal
- phylogenetic



Many of these models can get **very slow** even for moderate N .

→ **INLA** is designed for that. Very efficient !

Summary

Random or fixed effects ?

Classical / oldschool interpretation

You are interested in overall means / overall effects, but must account for non-independence in data?

→ **site as a random grouping factor**

~ 1 + (1|site) random intercepts

~ x + (1|site) random intercepts

~ x + (x|site) random intercepts & slopes

You are interested in differences in means / differences in effects between sites?

→ **site as fixed effect**

~ site ANOVA

~ x + site ANCOVA (diff. in intercepts)

~ x * site ANCOVA (diff. in intercepts & slopes)

Modern / relaxed interpretation

Can use random factor models for both to overcome difficulties like unbalanced / heterogeneous data

What's Bayesian about it ?


- Sometimes lme4 just does not converge !
- Correct quantification of uncertainties *on all levels*.
- Works for all numbers of groups or numbers of obs.
Already 2-level grouping factor is OK.
- Interpretation “easy” through levels of priors
- Frequentist assumption: Random effects drawn from a larger population of other, unmeasured groups
- Bayes: no such limitation
- Non-normal distribution of random effects?
→ Bayes (Stan) can handle that!



Summary

- Bayesian stats does not free you from model assumptions !
- If residuals are not independent, you must account for it.
- Otherwise: Fixed effects wrongly confident, CIs too narrow
- Thinking hard about your data will reveal random effects structure.
- Formula notation just as in lme4.
- Combine any kind of (generalized) linear (or nonlinear) model with discrete (random) grouping factors and/or continuous correlation structure ...
- ... in one R-package !!

The Bayesian 3D printer



Source: stockcake.com

Further reading

Bolker, B. M. (2015). Linear and generalized linear mixed models. In Fox, G. A. , Negrete-Yankelevich, S. and Sosa, V. J. (eds). *Ecological Statistics: Contemporary theory and application*. *Oxford Academic*.

<https://doi.org/10.1093/acprof:oso/9780199672547.003.0014>

Fieberg, J. (2024). Statistics 4 Ecologists. <https://statistics4ecologists-v2.netlify.app/> [Chapter 18]

Gelman, A., Hill, J. (2006). *Data Analysis Using Regression and Multilevel/Hierarchical Models*. Cambridge University Press

<http://www.stat.columbia.edu/~gelman/arm/>

Harrison, X. A., Donaldson, L., Correa-Cano, M. E., Evans, J., Fisher, D. N., Goodwin, C. E. D., Robinson, B. S., Hodgson, D. J., & Inger, R. (2018). A brief introduction to mixed effects modelling and multi-model inference in ecology. *PeerJ*, 6, e4794.

<https://doi.org/10.7717/peerj.4794>

Marshall, D. J. (2024). Principles of experimental design for ecology and evolution. *Ecology Letters*, 27(4), e14400.

<https://doi.org/10.1111/ele.14400>

McElreath, R. (2020). *Statistical Rethinking: A Bayesian Course with Examples in R and STAN* (2nd ed.). *Chapman and Hall/CRC*. <https://doi.org/10.1201/9780429029608> [Chapter 13]

Silk, M. J., Harrison, X. A., & Hodgson, D. J. (2020). Perils and pitfalls of mixed-effects regression models in biology. *PeerJ*, 8, e9522. <https://doi.org/10.7717/peerj.9522>

Warton, D. (2022). *Eco-Stats: Data Analysis in Ecology*. *Springer (Methods in Statistical Ecology)* [Chapter 6]