

Practical 5: Generalized linear models

Benjamin Rosenbaum

Table of contents

Poisson regression	2
Binomial regression	5
Overdispersion	13
Poisson	14
Negative Binomial	18
Distributional model	23
Logistic regression	29
Model 1	31
Model 2	34
Model 3	39

We learn how to use some classic GLMs, and a distributional model. We focus on prior specifications and appropriate posterior predictive checks. Sometimes we have to apply ggplot tricks for good conditional effects plots.

```
rm(list=ls())
library("brms")
library("ggplot2")
library("arm")
library("emmeans")
library("performance")
library("Data4Ecologists")
library("cowplot")
try(dev.off())
```

Poisson regression


We start with a simple comparison of 2 group means (frequentists would use a t-test), but the response are counts (number of slugs found per plot, 2 field types, from Dta4Ecologists package).

Deterministic part: $\mu = \mu(\text{field}) \rightarrow \mu_1$ or μ_2

Stochastic part: $\text{slugs} \sim \text{Poisson}(\mu)$

No log-link is used for this simple model, just “identity link”.

```
data(slugs)
ggplot(slugs, aes(field, slugs)) +
  geom_jitter(alpha=0.5, width=0.1, height=0.05)
```



It's always good to check brms default priors in a GLM:

```
default_prior(slugs ~ field,
  family = poisson(link=identity),
  data = slugs)
```

	prior	class	coef	group	resp	dpar	nlpar	lb	ub	source
	(flat)	b								default
	(flat)	b	fieldRookery							(vectorized)
student_t(3, 1, 2.5)		Intercept								default

We provide a vague prior for effect (difference in means, dummy-coded) and ensure that the intercept is positive with lower boundary (lb=0)

```
fit.slugs = brm(slugs ~ field,
  family = poisson(link=identity),
  prior =
    prior(normal(0,2), class=b) +
    prior(student_t(3, 1, 2.5), class=Intercept, lb=0),
  data = slugs)
```

Check convergence

```
summary(fit.slugs, prior=TRUE)
```

```
Family: poisson
Links: mu = identity
Formula: slugs ~ field
Data: slugs (Number of observations: 80)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000
```

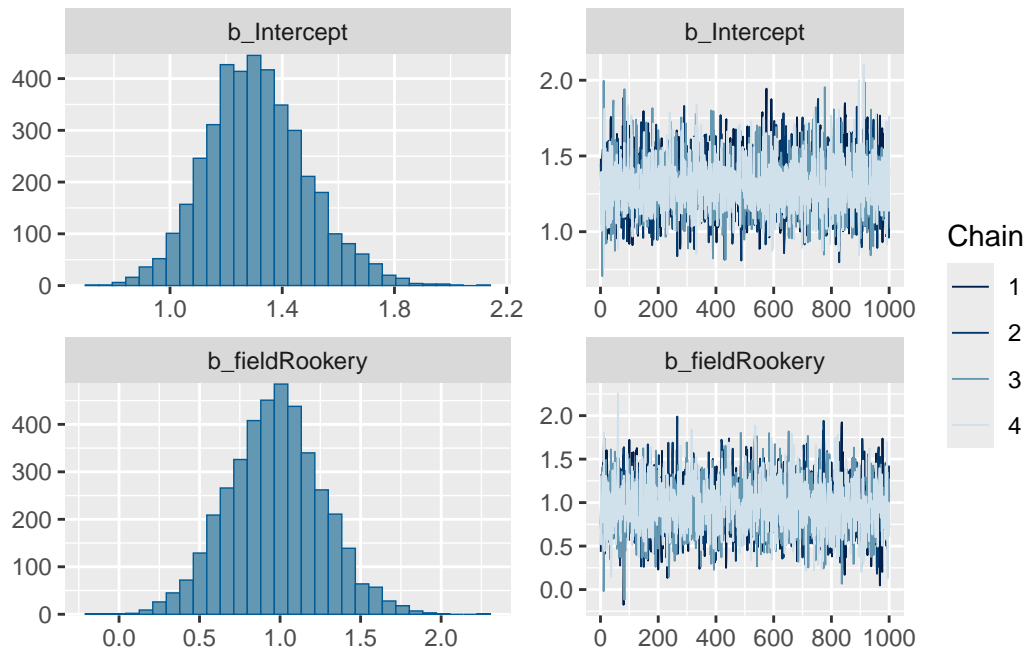
```
Priors:
b ~ normal(0, 2)
<lower=0> Intercept ~ student_t(3, 1, 2.5)
```

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	1.31	0.18	0.98	1.70	1.00	3355	2765
fieldRookery	0.97	0.29	0.40	1.58	1.00	2634	2123

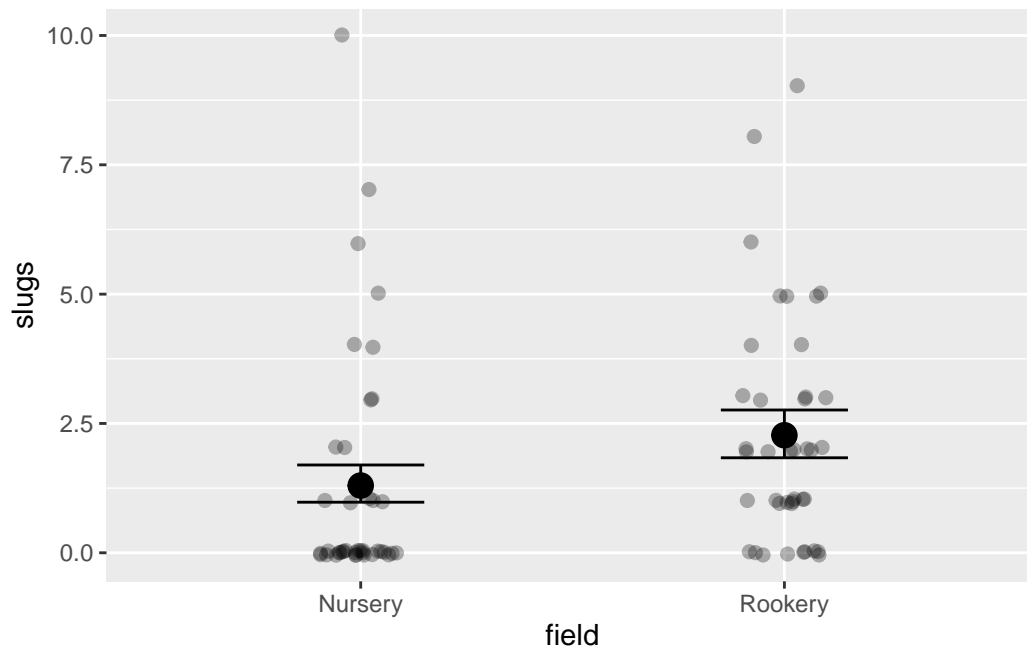
Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```
plot(fit.slugs)
```



Posterior predictions

```
plot(conditional_effects(fit.slugs),
     points=T,
     point_args=list(alpha=0.3, width=0.1, height=0.05)
)
```



In this simple example of comparing two group-means, the summary table already tells us that `field=Rookery` has on average +0.97 more slugs (95% CI [0.40,1.58]).

Importantly, we get quite a different effect size if we had just fitted an LM (effect=0.94, but 95% CI covering zero: [-0.05,1.90]). The Poisson model for count data is the correct model here, especially when observations are close or equal to zero!

```
fit.slugs.lm = brm(slugs ~ field,
  prior =
    prior(normal(0,2), class=b) +
    prior(student_t(3, 1, 2.5), class=Intercept, lb=0),
  data = slugs)
```

```
fixef(fit.slugs.lm) |> round(3)
```

	Estimate	Est.Error	Q2.5	Q97.5
Intercept	1.303	0.352	0.612	1.990
fieldRookery	0.935	0.494	-0.054	1.904

Binomial regression

This experiment is about mice infected with parasites of the genus *Cryptosporidium*. Each replicate has a certain number of mice inoculated by some Dose of parasite oocysts. Then the number of infected mice is counted. (from Quian (2016) Environmental and Ecological Statistics with R)


Fit a dose-response model.

Question: What parasite dose is required to infect 75% of a mouse population?

Deterministic part: $\text{logit}(\mu) = a + b \cdot \text{Dose}$

Stochastic part: $Y_{\text{infected}} \sim \text{Binomial}(Y_{\text{inoculated}}, \mu)$

```
df = read.csv("https://raw.githubusercontent.com/songsqian/eesR/refs/heads/master/R/Data/cryp
  sep=" ")
df.sub = subset(df, Source=="Finch")
df.sub$Y = as.integer(df.sub$Y)
ggplot(df.sub, aes(Dose, Y/N)) + geom_point(alpha=0.5)
```



Again, we want to scale the predictor `Dose`. Here, we directly write `scale(Dose)` in the model formula. The advantage is that priors and parameters on z-scale, but plots / conditional_effects are on original scale.

```
default_prior(Y | trials(N) ~ scale(Dose),
              family = binomial(link=logit),
              data = df.sub )
```

	prior	class	coef	group	resp	dpar	nlpar	lb	ub	source
	(flat)	b								default
	(flat)	b	scaleDose							(vectorized)
student_t(3, 0, 2.5)		Intercept								default

We add a vaguely informative prior on the (linear scale) slope. The effect is expected to be positive.

```
fit.crypto = brm(Y | trials(N) ~ scale(Dose),
                 family = binomial(link=logit),
                 prior = prior(normal(1,1), class=b),
                 data = df.sub )
```

Check convergence

```
summary(fit.crypto, prior=T)
```

```
Family: binomial
Links: mu = logit
Formula: Y | trials(N) ~ scale(Dose)
Data: df.sub (Number of observations: 43)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000
```

Priors:

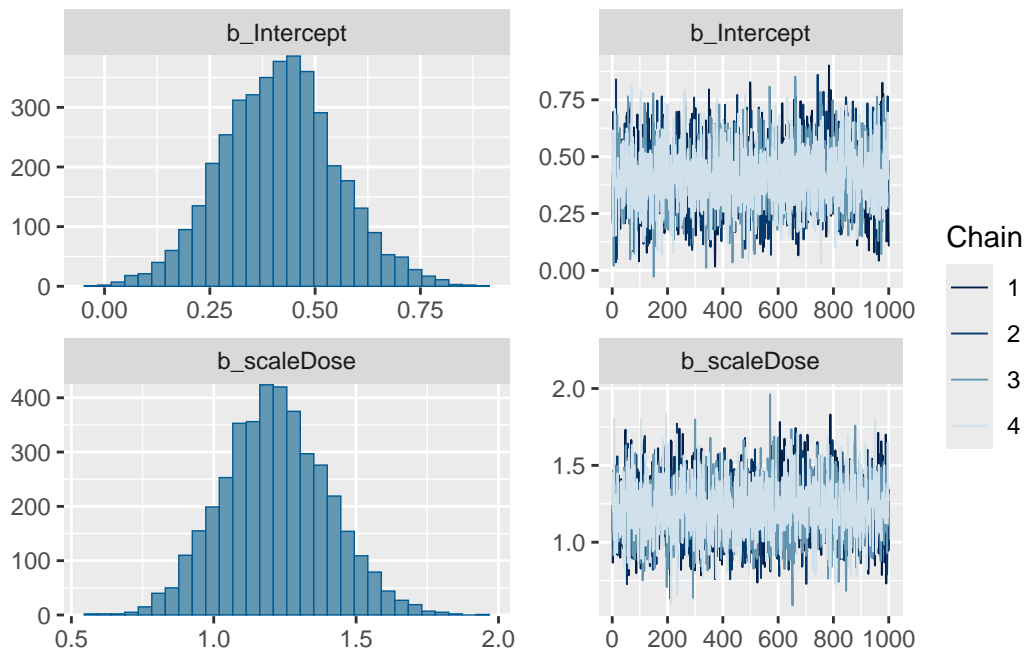
```
b ~ normal(1, 1)
Intercept ~ student_t(3, 0, 2.5)
```

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	0.41	0.14	0.15	0.69	1.00	3019	2729
scaleDose	1.22	0.19	0.86	1.59	1.00	2677	2494

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).


```
plot(fit.crypto)
```



`conditional_effects()` plots predictions for $N=1$ trials only. `points=TRUE` just adds Y values, but we need Y/N values!


```
plot(conditional_effects(fit.crypto, effects="Dose"), points=TRUE)
```

Setting all 'trials' variables to 1 by default if not specified otherwise.




Some ggplot-tricks to the rescue!! Note that `conditional_effects()` now plots everything on the predictor's original scale, not on the z-scale.

```
p1 = plot(conditional_effects(fit.crypto, effects="Dose"),  
          points=FALSE, plot=FALSE)  
p1[[1]] + geom_point(data=df.sub,  
                     aes(Dose, Y/N),  
                     alpha=0.6, size=1.5,  
                     inherit.aes=FALSE)
```





Posterior predictive checks are not very pretty, but we didn't measure any other predictors to improve it.

```
p1 = pp_check(fit.crypto, ndraws=100)
p2 = pp_check(fit.crypto, type="scatter_avg")
plot_grid(p1,p2)
```




```
fitted = fitted(fit.crypto)
residuals = residuals(fit.crypto)
binnedplot(fitted, residuals)
```



What parasite dose is needed to get 75% of mice infected?

```
p1 = plot(conditional_effects(fit.crypto, effects="Dose"),
          points=FALSE, plot=FALSE)
p1[[1]] + geom_hline(yintercept=0.75, linetype=2, col="red")
```




For Dose=191, the mean regression curve crosses the 75%-infected line, **on average** 75% are infected. But this means that we are **only 50% certain** that at least 75% are infected, $P(\mu > 0.75) \approx 0.5$.

```
fitted(fit.crypto,
       newdata = data.frame(Dose=191, N=1))
```

```
      Estimate Est.Error      Q2.5      Q97.5
[1,] 0.7513291 0.03655965 0.6768719 0.8215464
```

```
fitted.191 = fitted(fit.crypto,
                    newdata = data.frame(Dose=191, N=1),
                    summary=FALSE)
hist(fitted.191)
abline(v=0.75, col="red", lwd=2)
```

Histogram of fitted.191

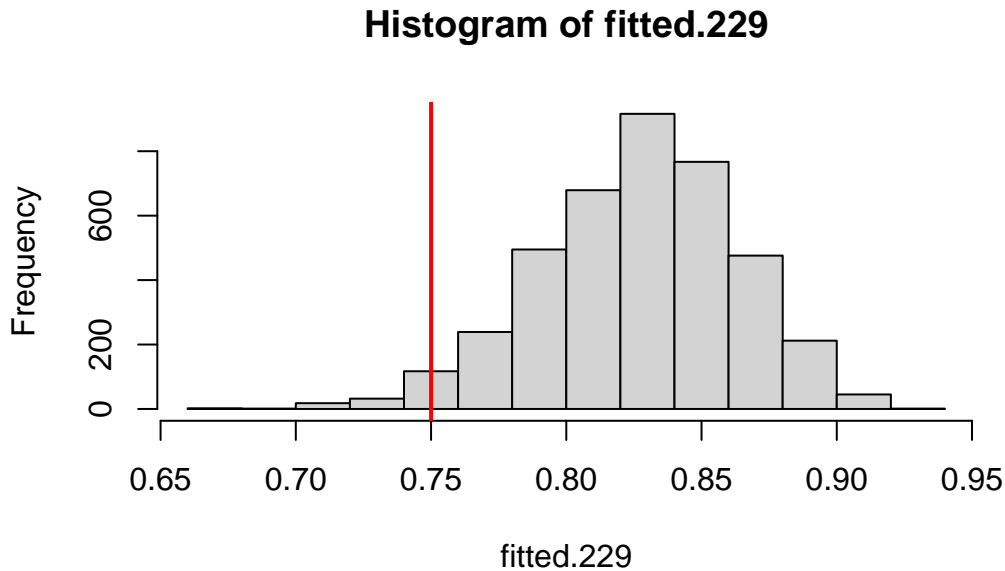


Rather, we'd have to supply a dose where the 75%-line crosses the fitted lower quantile (Q2.5), which happens around Dose=229. Then we are 97.5% certain that at least 75% of mice are infected, $P(\mu > 0.75) = 0.975$.

```
fitted(fit.crypto,
       newdata=data.frame(Dose=229, N=1),
       probs=0.025)
```

```
      Estimate Est.Error      Q2.5
[1,] 0.8269148 0.03625958 0.7504071
```


```
fitted.229 = fitted(fit.crypto,
                    newdata = data.frame(Dose=229, N=1),
                    summary=FALSE)
hist(fitted.229)
abline(v=0.75, col="red", lwd=2)
```



Overdispersion

We model the daily growth of young chicken (from datasets package). We use only the early exponential growth phase (looks linear on the logscale). Later, growth slows down and a the nonlinear von Bertalanffy growth function is more appropriate.

```
data(ChickWeight)
p1 = ggplot(ChickWeight, aes(Time, weight)) +
  geom_jitter(alpha=0.5, width=0.1, height=0.05)
p2 = ggplot(ChickWeight, aes(Time, log(weight))) +
  geom_jitter(alpha=0.5, width=0.1, height=0.05)
plot_grid(p1,p2)
```



```
ChickWeight = subset(ChickWeight, Time<15)
```

Poisson

The response weight is measured in integers, so we could use Poisson here. Using a log-link, we fit a linear model.

Deterministic part: $\log(\mu) = a + b \cdot \text{time}$
 Stochastic part: $\text{weight} \sim \text{Poisson}(\mu)$

```
default_prior(weight ~ Time,
              family = poisson(link=log),
              data = ChickWeight)
```

	prior	class	coef	group	resp	dpar	nlpar	lb	ub	source
	(flat)		b							default
	(flat)		b Time							(vectorized)
student_t(3, 4.3, 2.5)		Intercept								default

We add a weak prior for the slope, but growth must be positive (chickens don't shrink).

```
fit.growth.1 = brm(weight ~ Time,  
  family = poisson(link=log),  
  prior = prior(normal(0,1), class=b, lb=0),  
  data = ChickWeight)
```

Check convergence

```
summary(fit.growth.1)
```


```
Family: poisson  
Links: mu = log  
Formula: weight ~ Time  
Data: ChickWeight (Number of observations: 393)  
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;  
total post-warmup draws = 4000
```

Regression Coefficients:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	3.75	0.01	3.72	3.77	1.00	1994	1903
Time	0.09	0.00	0.09	0.09	1.00	2180	1779



Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```
plot(fit.growth.1)
```




Posterior predictions for fitted values (deterministic part) look OK, but the predicted values don't correspond to that variance in the data (95% prediction intervals should contain ~95% of the data)

```
p1 = plot(conditional_effects(fit.growth.1),
          points=T,
          point_args=list(alpha=0.25),
          plot=F)
p2 = plot(conditional_effects(fit.growth.1, method="posterior_predict"),
          points=T,
          point_args=list(alpha=0.25),
          plot=F)
plot_grid(p1[[1]], p2[[1]])
```

```
p1 = pp_check(fit.growth.1, ndraws=100)
p2 = pp_check(fit.growth.1, type="scatter_avg")
plot_grid(p1,p2)
```



Negative Binomial

Deterministic part: $\log(\mu) = a + b \cdot \text{time}$

Stochastic part: $\text{weight} \sim \text{NegativeBinomial}(\mu, \phi)$

ϕ is an additional shape parameter modeling overdispersion, meaning the variance of the response does not scale linearly with the response itself (as it does in Poisson).

```
default_prior(weight ~ Time,
               family = negbinomial(link=log),
               data = ChickWeight)
```

	prior	class	coef	group	resp	dpar	nlpar	lb	ub	source
	(flat)		b							default
	(flat)		b	Time						(vectorized)
student_t(3, 4.3, 2.5)		Intercept								default
inv_gamma(0.4, 0.3)		shape						0		default

Again, we add a weak but strictly positive prior on the slope, brms provides default values also for ϕ

```
fit.growth.2 = brm(weight ~ Time,
  family = negbinomial(link=log),
  prior = prior(normal(0,1), class=b, lb=0),
  data = ChickWeight)
```

Check convergence

```
summary(fit.growth.2)
```

```
Family: negbinomial
Links: mu = log; shape = identity
Formula: weight ~ Time
Data: ChickWeight (Number of observations: 393)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000
```

Regression Coefficients:

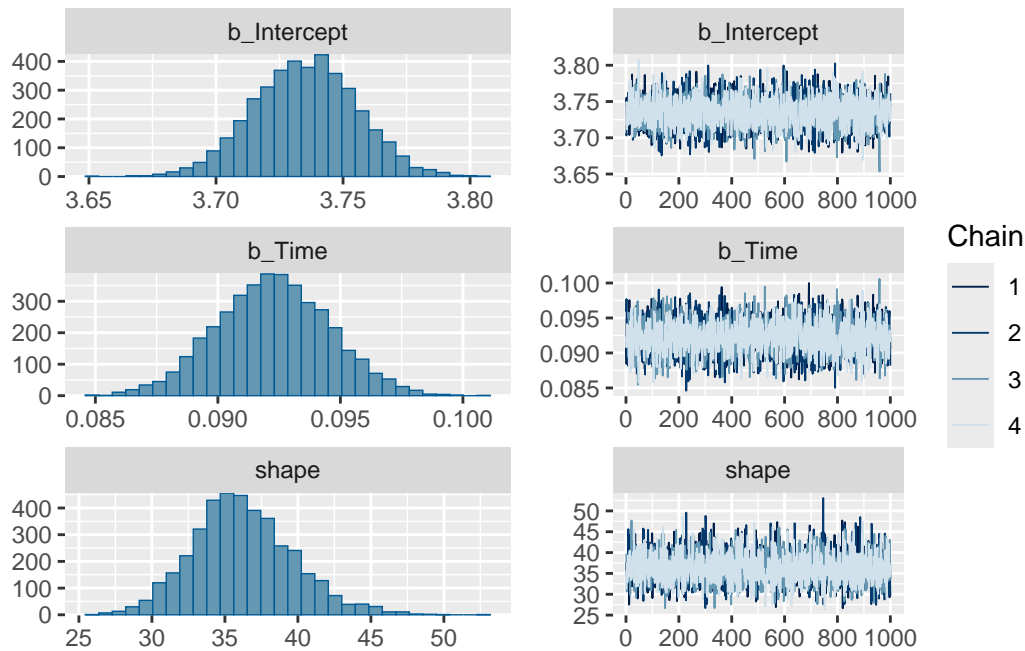
	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	3.74	0.02	3.70	3.77	1.00	4080	2769
Time	0.09	0.00	0.09	0.10	1.00	4537	3190

Further Distributional Parameters:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
shape	36.25	3.46	30.03	43.87	1.00	4312	2861



Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```
plot(fit.growth.2)
```




Looking at the same plot as before, the negative binomial model is more appropriate and replicates the variance of the observations much better.

```
p1 = plot(conditional_effects(fit.growth.2),
          points=T,
          point_args=list(alpha=0.25),
          plot=F)
p2 = plot(conditional_effects(fit.growth.2, method="posterior_predict"),
          points=T,
          point_args=list(alpha=0.25),
          plot=F)
plot_grid(p1[[1]], p2[[1]])
```



```
p1 = pp_check(fit.growth.2, ndraws=100)
p2 = pp_check(fit.growth.2, type="scatter_avg")
plot_grid(p1,p2)
```



Finally, model comparison reveals that the negative binomial model fits the data better.

```
L00(fit.growth.1, fit.growth.2)
```

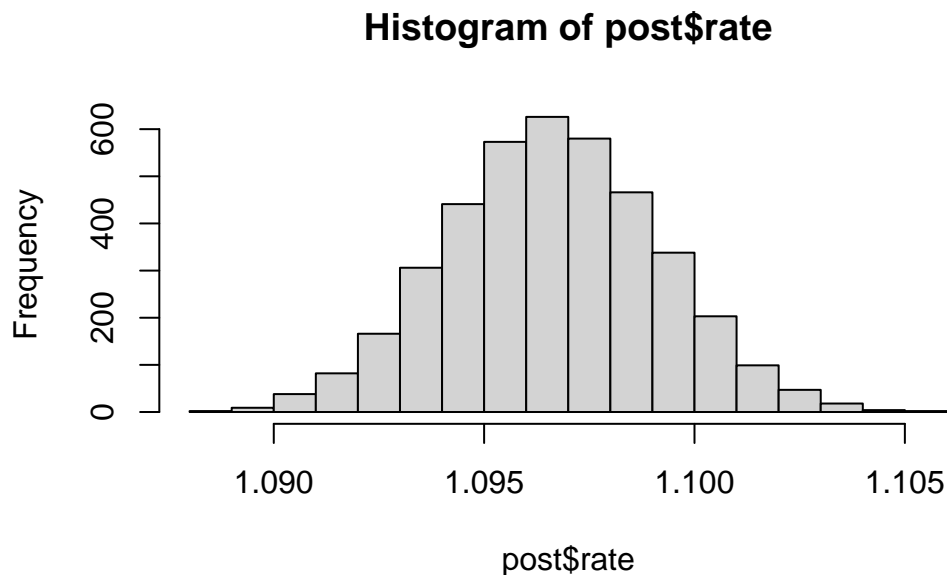
	elpd_diff	se_diff
fit.growth.2	0.0	0.0
fit.growth.1	-327.9	66.2

Question: What is the average daily growth rate?

$$\log(\mu) = a + b \cdot \text{time} \quad \rightarrow \quad \mu = \exp(a) \cdot \exp(b)^{\text{time}}$$

$\exp(b)$ is average daily growth rate, but we must use the whole posterior distribution! Don't just use mean \bar{b} and compute $\exp(\bar{b})$. Jensen's inequality says you might get a biased value with nonlinear parameter transformations!

```
post = as_draws_df(fit.growth.2)
post$rate = exp(post$b_Time)
hist(post$rate)
```



```
mean(post$rate)
```

```
[1] 1.096641
```

```
quantile(post$rate, prob=c(0.05, 0.95))
```

```
      5%      95%
1.092519 1.100806
```

The daily growth rate is 1.096, this means an 9.6% increase every day.

Distributional model

We use the same data, but treat the response as a continuous variable. We want to perform a simple linear regression on $\log(\text{weight})$ first.

Deterministic part: $\mu = a + b \cdot \text{time}$

Stochastic part: $\log(\text{weight}) \sim \text{Normal}(\mu, \sigma)$


```
fit.growth.3 = brm(log(weight) ~ Time,
  prior = prior(normal(0,1), class=b),
  data = ChickWeight)
```

Again, we see that prediction intervals don't represent the variance of data well.

```

p1 = plot(conditional_effects(fit.growth.3),
          points=T,
          point_args=list(alpha=0.25),
          plot=F)
p2 = plot(conditional_effects(fit.growth.3, method="posterior_predict"),
          points=T,
          point_args=list(alpha=0.25),
          plot=F)
plot_grid(p1[[1]],p2[[1]])


```



```

p1 = pp_check(fit.growth.3, ndraws=100)
p2 = pp_check(fit.growth.3, type="scatter_avg")
plot_grid(p1,p2)

```





Linear model assumptions are not satisfied

```
check_model(fit.growth.3, check=c("linearity", "homogeneity", "qq", "normality"))
```


Linearity

Reference line should be flat and horizontal




Homogeneity of Variance

Reference line should be flat and horizontal




Normality of Residuals

Points should fall along the line



Normality of Residuals

Distribution should be close to the normal distribution



We introduce a **distributional model**, where not only μ , but also σ is a function of the predictor *time*.

Deterministic part: $\mu = a + b \cdot \text{time}$

$\log(\sigma) = c + d \cdot \text{time}$ (log-link to keep $\sigma > 0$)

Stochastic part: $\log(\text{weight}) \sim \text{Normal}(\mu, \sigma)$

Both model parts for μ and σ are combined with `bf()` (short for `brmsformula`). For sigma, the log-link is provided automatically

```
default_prior(bf(log(weight) ~ Time,
                 sigma ~ Time),
              data = ChickWeight)
```

	prior	class	coef	group	resp	dpar	nlpar	lb	ub	source
	(flat)		b							default
	(flat)		b	Time						(vectorized)
student_t(3, 4.3, 2.5)		Intercept								default
	(flat)		b			sigma				default
	(flat)		b	Time		sigma				(vectorized)
student_t(3, 0, 2.5)		Intercept				sigma				default

We provide weak priors for both effects (slope for μ and σ with time)

```
fit.growth.4 = brm(bf(log(weight) ~ Time,
                     sigma ~ Time),
                  prior =
                    prior(normal(0,1), class=b, coef=Time) +
                    prior(normal(0,1), class=b, dpar=sigma),
                  data = ChickWeight)
```

Check convergence

```
summary(fit.growth.4, prior=T)
```

```
Family: gaussian
Links: mu = identity; sigma = log
Formula: log(weight) ~ Time
         sigma ~ Time
Data: ChickWeight (Number of observations: 393)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
```

```
total post-warmup draws = 4000
```

Priors:

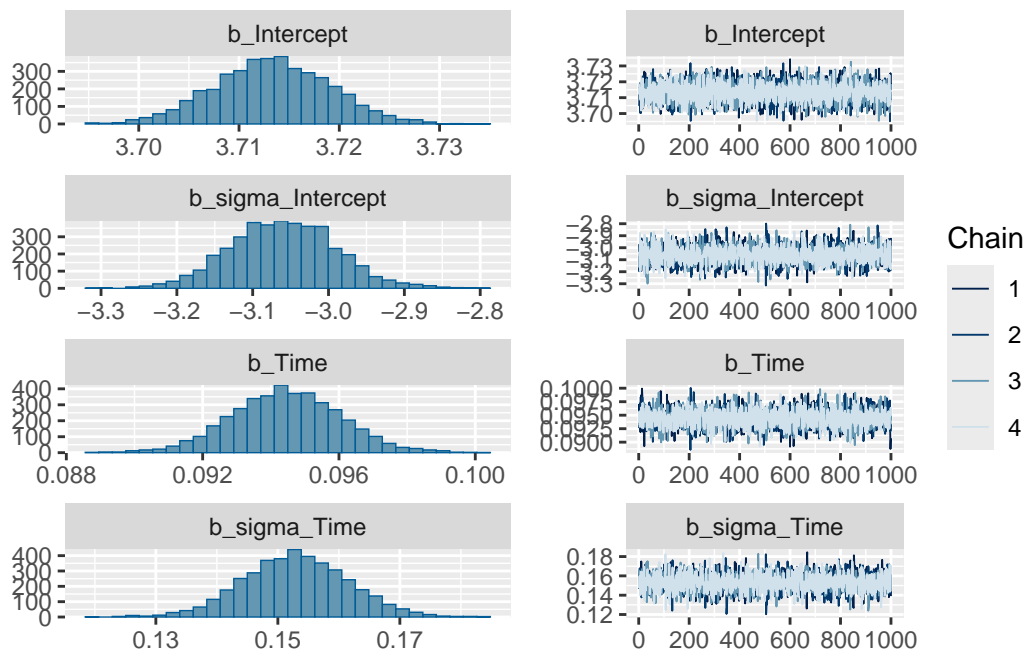
```
b_Time ~ normal(0, 1)
b_sigma ~ normal(0, 1)
Intercept ~ student_t(3, 4.3, 2.5)
Intercept_sigma ~ student_t(3, 0, 2.5)
```

Regression Coefficients:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	3.71	0.01	3.70	3.72	1.00	4494	3323
sigma_Intercept	-3.06	0.07	-3.19	-2.92	1.00	2997	2371
Time	0.09	0.00	0.09	0.10	1.00	3013	2094
sigma_Time	0.15	0.01	0.14	0.17	1.00	3155	2368

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```
plot(fit.growth.4)
```




Posterior predictions are looking much better now. We don't use the `check_model()` here, because it's not a linear model anymore.

```

p1 = plot(conditional_effects(fit.growth.4),
          points=T,
          point_args=list(alpha=0.25),
          plot=F)
p2 = plot(conditional_effects(fit.growth.4, method="posterior_predict"),
          points=T,
          point_args=list(alpha=0.25),
          plot=F)
plot_grid(p1[[1]],p2[[1]])


```



```

p1 = pp_check(fit.growth.4, ndraws=100)
p2 = pp_check(fit.growth.4, type="scatter_avg")
plot_grid(p1,p2)

```



Model comparison reveals that the distributional model is preferred.

```
LOO(fit.growth.3, fit.growth.4)
```

	elpd_diff	se_diff
fit.growth.4	0.0	0.0
fit.growth.3	-131.1	22.3

Attention: Do not use LOOIC / WAIC / AIC / BIC / DIC for comparing models for continuous responses (Normal, Lognormal, Gamma, Beta, ...) to models for discrete responses (Binomial, Poisson, Negative Binomial, ...)

Continuous distributions are based on probability density functions, while discrete distributions have probability mass functions, which you can't compare.

Logistic regression

From Qian, S. (2016) Environmental and Ecological Statistics with R

This field experiment is about seed predation by rodents (response *Predation* = {0,1}). Seeds were placed in gauze bags in 4 topographic locations (predictor *topo*). Sites were visited in 6 sampling campaigns (predictor *time*). Also control for continuous predictor *seed.weight*.

Deterministic part: $\text{logit}(\mu) = \dots$

Stochastic part: $Predation \sim \text{Bernoulli}(\mu)$

Start with `Predation~log(seed.weight)`

Test if there is a difference between locations (`+topo`)


Also use time as factorial predictor (`+time`)

Question: Does predation rate reach a stable value in the last three time periods?

```
df = read.csv("https://raw.githubusercontent.com/songsqian/eesR/refs/heads/master/R/Data/seed.csv",
              sep=",")
df = subset(df, Predation %in% 0:1)
df$Predation = as.integer(df$Predation)
df$time = as.factor(df$time)
df$topo = as.factor(df$topo)
df$seed = scale(log(df$seed.weight))
head(df)
```

	species	seed.weight	time	topo	ground	Predation	seed
1	8	4250.0	3	1	2	0	2.10606568
2	5	28.5	1	1	2	0	0.01142359
3	5	28.5	3	1	2	0	0.01142359
4	2	4.5	1	1	2	0	-0.76110864
5	2	4.5	2	1	2	0	-0.76110864
6	2	4.5	3	1	2	0	-0.76110864

```
ggplot(df, aes(log(seed.weight), Predation)) +
  geom_jitter(height=0.05, width=0.05, alpha=0.2)
```



Model 1

We start with a simple logistic regression, just with the continuous predictor `seed`. Note that priors for intercept and slope are on linear scale (η in the lecture), not on response scale (μ in the lecture).

```
df.seed.1 = brm(Predation ~ seed,
  family = bernoulli(link=logit),
  prior = prior(normal(1,1), class=b),
  data = df )
```

Check convergence

```
summary(df.seed.1, prior=T)
```

```
Family: bernoulli
Links: mu = logit
Formula: Predation ~ seed
Data: df (Number of observations: 1142)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000
```

Priors:

$b \sim \text{normal}(1, 1)$

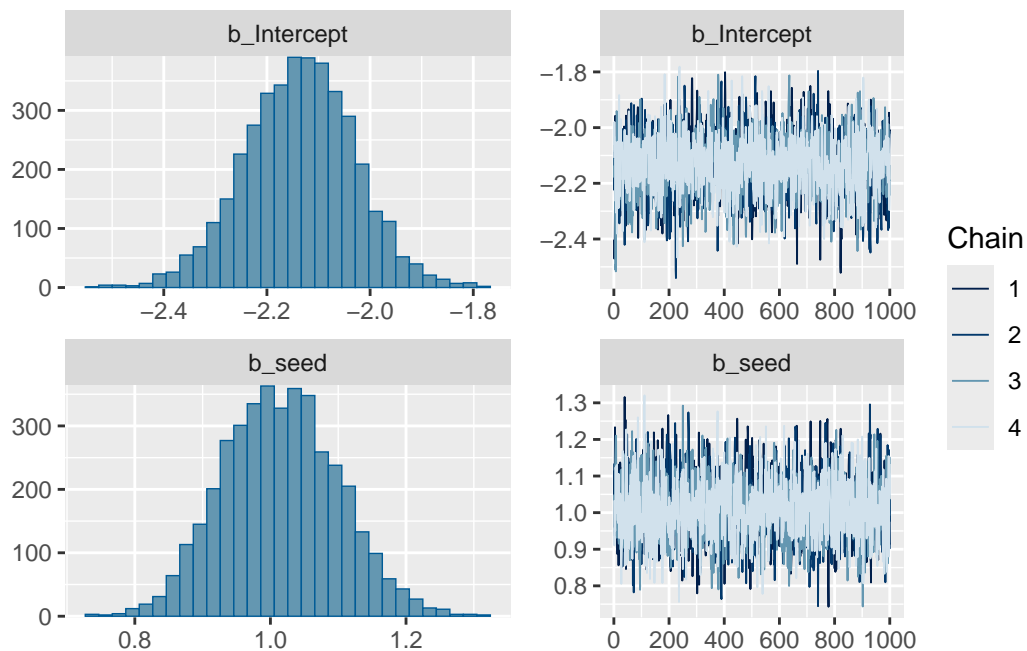
$\text{Intercept} \sim \text{student_t}(3, 0, 2.5)$

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	-2.14	0.11	-2.35	-1.93	1.00	1743	2201
seed	1.02	0.09	0.86	1.19	1.00	1876	2271


Draws were sampled using `sampling(NUTS)`. For each parameter, `Bulk_ESS` and `Tail_ESS` are effective sample size measures, and `Rhat` is the potential scale reduction factor on split chains (at convergence, `Rhat` = 1).

```
plot(df.seed.1)
```




Posterior predictive checks

```
plot(conditional_effects(df.seed.1),  
     points = T,  
     point_args = list(alpha=0.1, width=0.02, height=0.05)  
)
```

We use the binned residuals plot, which is appropriate for 0/1 responses. It doesn't look great, but we will add more predictors

```
fitted = fitted(df.seed.1)
residuals = residuals(df.seed.1)
binnedplot(fitted, residuals)
```



Model 2

Now we add the categorical predictor `topo` to account for differences in location.

```
default_prior(Predation ~ seed + topo,
              family = bernoulli(link=logit),
              data = df )
```

	prior	class	coef	group	resp	dpar	nlpar	lb	ub	source
	(flat)	b								default
	(flat)	b	seed							(vectorized)
	(flat)	b	topo2							(vectorized)
	(flat)	b	topo3							(vectorized)
	(flat)	b	topo4							(vectorized)
student_t(3, 0, 2.5)		Intercept								default

We assign priors for seed slope (continuous) and topo-differences (categorical)

```
df.seed.2 = brm(Predation ~ seed + topo,
                 family = bernoulli(link=logit),
                 prior =
```

```
prior(normal(0,2), class=b) + # for all effects
prior(normal(1,1), class=b, coef=seed), # just for seed
data = df )
```

Check convergence

```
summary(df.seed.2, prior=T)
```

```
Family: bernoulli
Links: mu = logit
Formula: Predation ~ seed + topo
Data: df (Number of observations: 1142)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000
```

Priors:

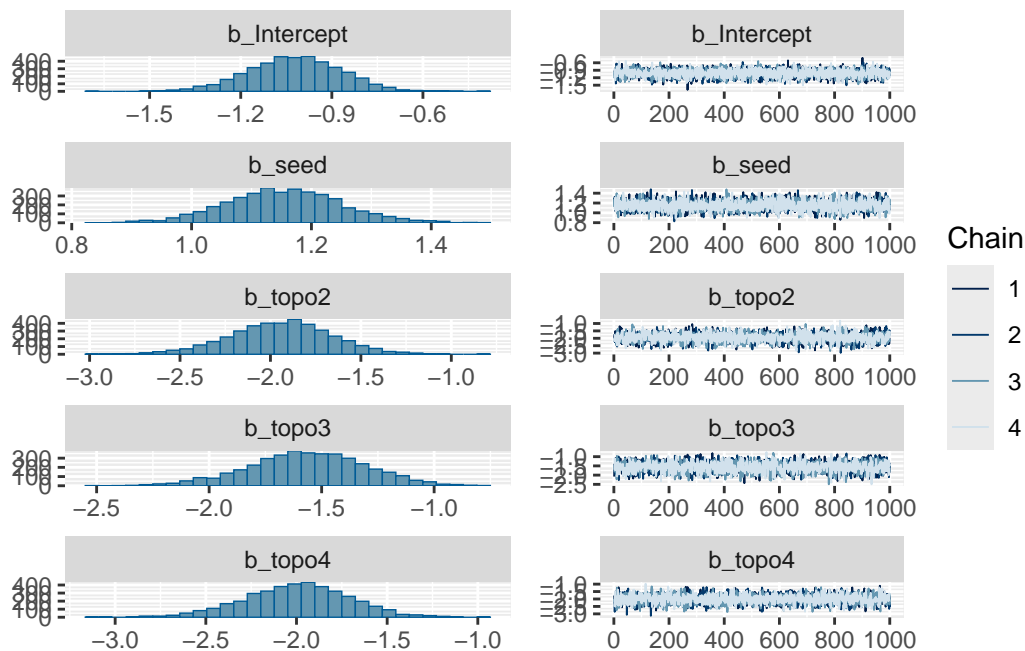
```
b ~ normal(0, 2)
b_seed ~ normal(1, 1)
Intercept ~ student_t(3, 0, 2.5)
```

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	-1.02	0.15	-1.32	-0.74	1.00	4301	3184
seed	1.15	0.10	0.97	1.35	1.00	2927	2778
topo2	-1.94	0.28	-2.52	-1.39	1.00	2927	2120
topo3	-1.57	0.26	-2.09	-1.07	1.00	3231	3086
topo4	-2.00	0.29	-2.58	-1.45	1.00	3392	2937

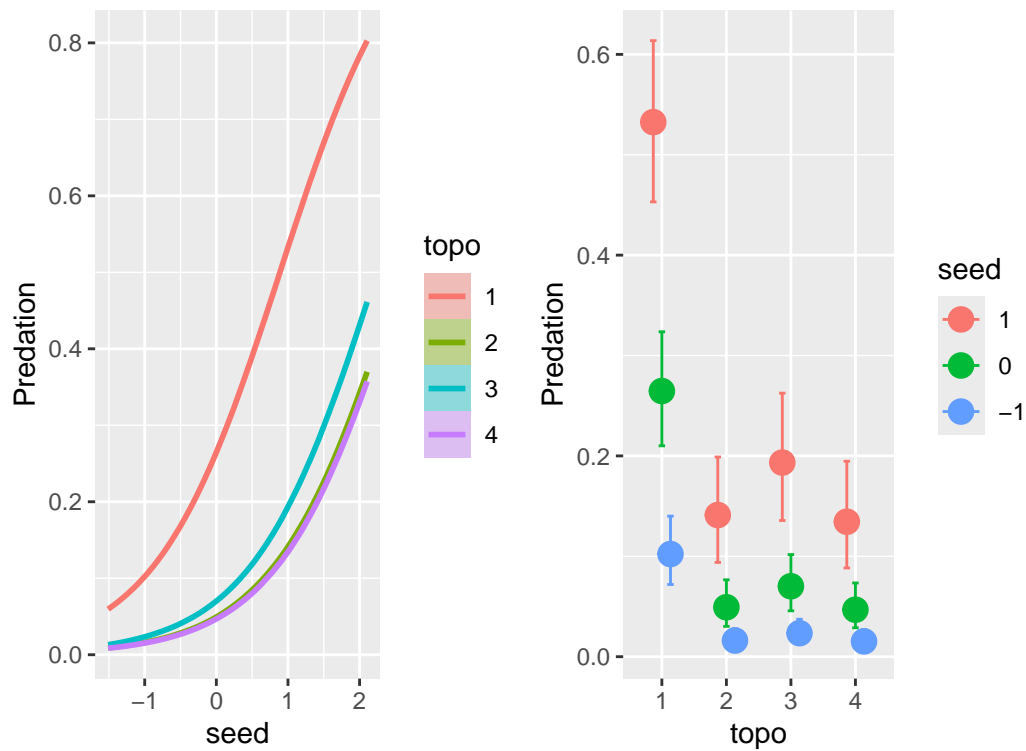
Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```
plot(df.seed.2)
```



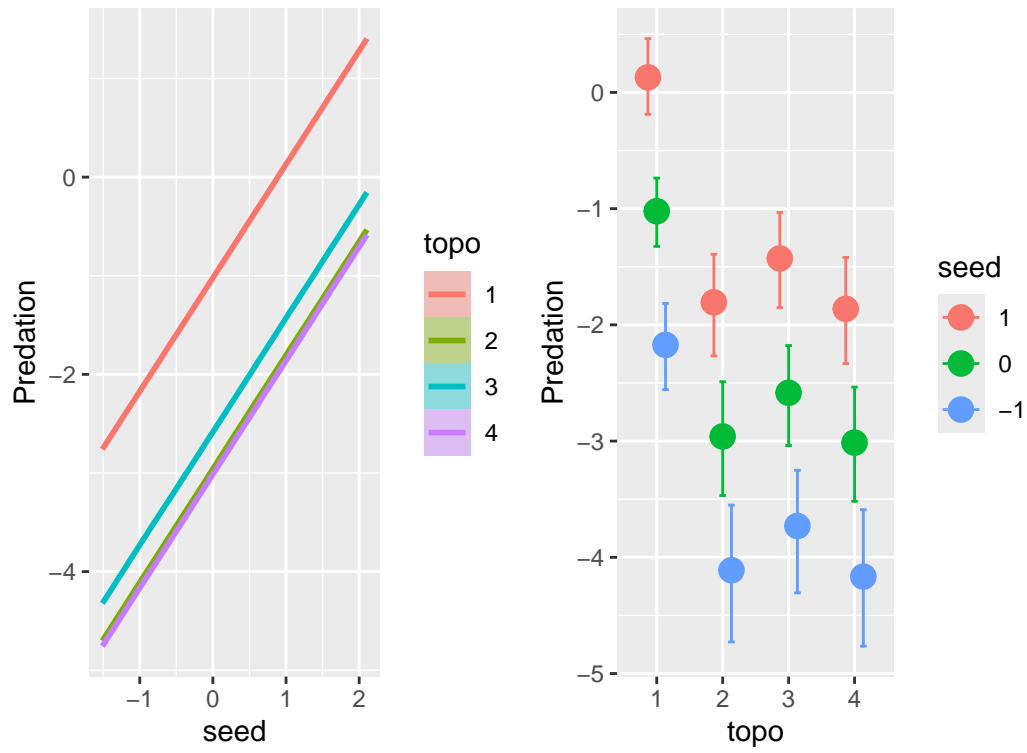
We see there is substantial variation between locations (topo)

```
p1 = plot(conditional_effects(df.seed.2,
                             effect="seed:topo",
                             prob=0),
          plot=F)
p2 = plot(conditional_effects(df.seed.2,
                             effect="topo:seed"),
          plot=F)
plot_grid(p1[[1]], p2[[1]])
```




This `~seed+topo` (generalized) ANCOVA without interaction produces parallel fitted values (intercept varies with `topo`) on the linear scale (`method="posterior_linpred"`).

```
p1 = plot(conditional_effects(df.seed.2,
                             effect="seed:topo",
                             prob=0,
                             method="posterior_linpred"),
          plot=F)
p2 = plot(conditional_effects(df.seed.2,
                             effect="topo:seed",
                             method="posterior_linpred"),
          plot=F)
plot_grid(p1[[1]], p2[[1]])
```



```
fitted = fitted(df.seed.2)
residuals = residuals(df.seed.2)
binnedplot(fitted, residuals)
```



Residuals still not looking great, but model comparison suggests that including `topo` improves the model.

```
L00(df.seed.1, df.seed.2)
```

	elpd_diff	se_diff
df.seed.2	0.0	0.0
df.seed.1	-40.3	9.3

Model 3

Now we add sampling campaign (`time`) as a factorial predictor.

```
df.seed.3 = brm(Predation ~ seed + topo + time,
  family = bernoulli(link=logit),
  prior =
    prior(normal(0,2), class=b) + # for all effects
    prior(normal(1,1), class=b, coef=seed), # just for seed
  data = df )
```

Check convergence

```
summary(df.seed.3, prior=T)
```

```
Family: bernoulli
Links: mu = logit
Formula: Predation ~ seed + topo + time
Data: df (Number of observations: 1142)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000
```

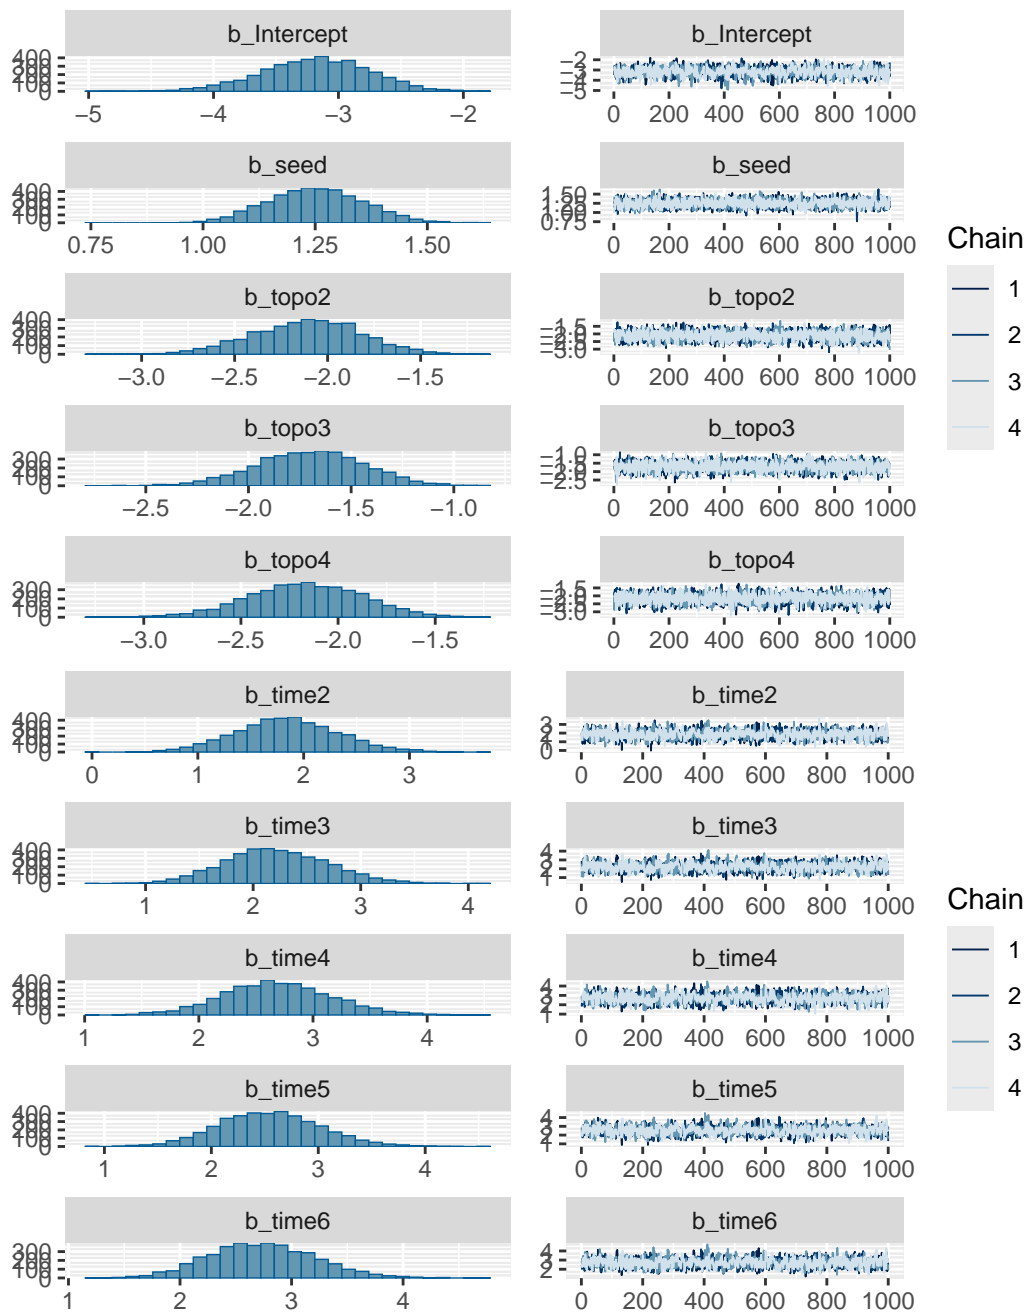
```
Priors:
b ~ normal(0, 2)
b_seed ~ normal(1, 1)
Intercept ~ student_t(3, 0, 2.5)
```

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	-3.19	0.43	-4.07	-2.39	1.00	1418	1671
seed	1.25	0.10	1.05	1.45	1.00	2017	2646
topo2	-2.11	0.29	-2.69	-1.56	1.00	2774	2972
topo3	-1.70	0.26	-2.23	-1.20	1.00	3010	3333
topo4	-2.17	0.29	-2.75	-1.61	1.00	2671	2714
time2	1.85	0.49	0.92	2.84	1.00	1642	1785
time3	2.22	0.48	1.32	3.19	1.00	1594	1777
time4	2.69	0.48	1.78	3.66	1.00	1497	1952
time5	2.54	0.48	1.65	3.49	1.00	1550	1941
time6	2.74	0.48	1.85	3.70	1.00	1541	2034

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```
plot(df.seed.3)
```


With 3 predictors, plotting conditional effects is getting more complicated, but the function can handle it.

The same plot as in Model 2, but when we don't specify the 3rd predictor, fitted values are shown at its reference level (`time=1`)

```

p1 = plot(conditional_effects(df.seed.3,
                             effect="seed:topo", # time at reference level
                             prob=0),
          plot=F)
p2 = plot(conditional_effects(df.seed.3,
                             effect="topo:seed"), # time at reference level
          plot=F)
plot_grid(p1[[1]],p2[[1]])

```

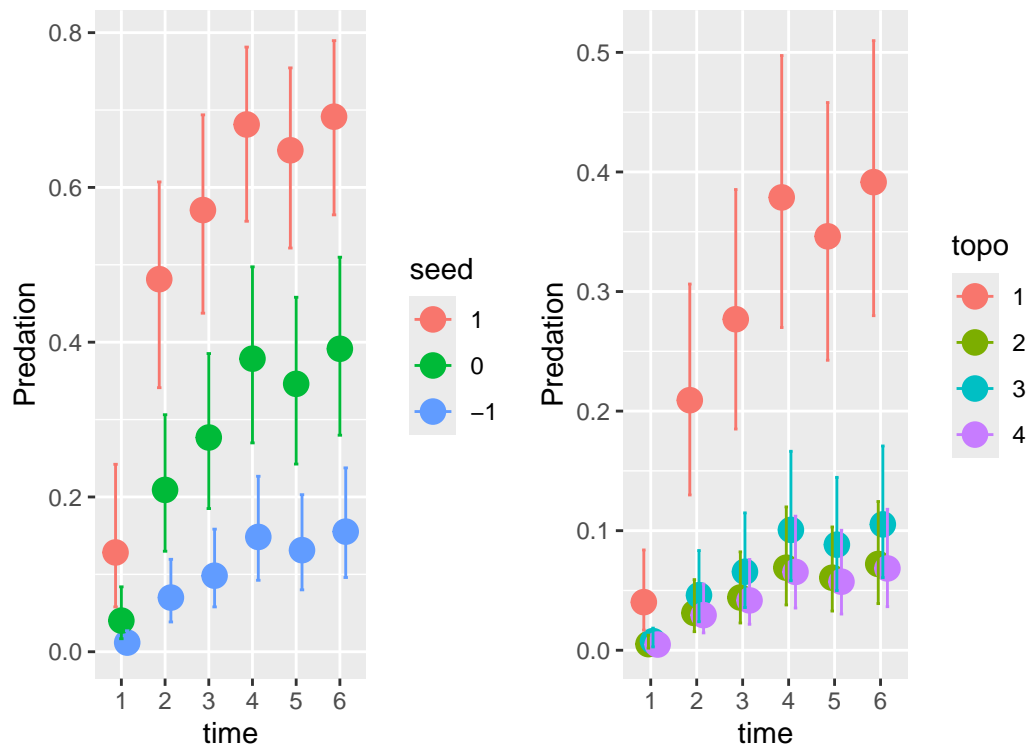


We're actually interested in prediction for different time-levels. Plotting time:seed and time:topo, the omitted 3rd predictor is held at its reference level.

```

p1 = plot(conditional_effects(df.seed.3,
                             effect="time:seed"), # topo at reference level
          plot=F)
p2 = plot(conditional_effects(df.seed.3,
                             effect="time:topo"), # seed at mean=0
          plot=F)
plot_grid(p1[[1]],p2[[1]])

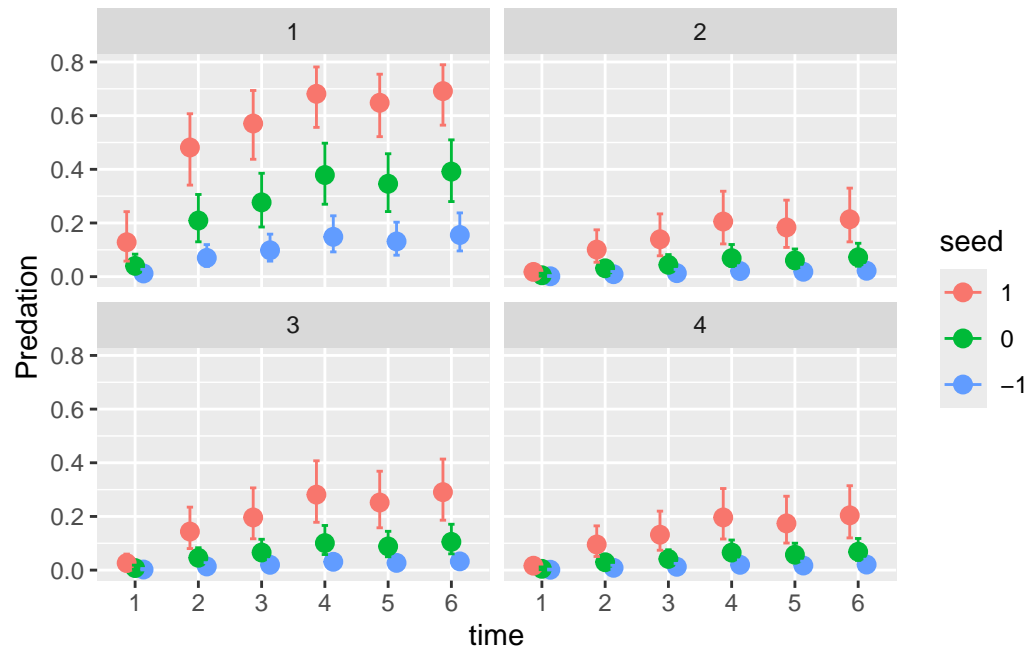
```




It already looks like predation rate is pretty constant in the last 3 campaigns.

By specifying `conditions=`, we can redo these plots for all levels of the 3rd predictor.

```
plot(conditional_effects(df.seed.3,
  effect="time:seed",
  conditions = make_conditions(df.seed.3,vars=c("topo"))))
```




```
plot(conditional_effects(df.seed.3,
  effect="time:topo",
  conditions = make_conditions(df.seed.3,vars=c("seed"))))
```



```
fitted = fitted(df.seed.3)
residuals = residuals(df.seed.3)
binnedplot(fitted, residuals)
```

Binned residual plot



Model comparison reveals that predation rate changes over time.

```
L00(df.seed.1, df.seed.2, df.seed.3)
```

	elpd_diff	se_diff
df.seed.3	0.0	0.0
df.seed.2	-28.4	5.3
df.seed.1	-68.7	10.8

But it still doesn't answer the question if predation rate reached a stable limit in the last 3 campaigns.

With emmeans, we can extract mean fitted values versus `time`, averaged over the remaining predictors. By default, the output is on linear scale, but it can also be displayed on response scale

```
emmeans(df.seed.3, ~time)
```

time	emmean	lower.HPD	upper.HPD
1	-4.66	-5.57	-3.86
2	-2.83	-3.37	-2.30
3	-2.46	-2.94	-1.95
4	-1.99	-2.43	-1.53
5	-2.13	-2.60	-1.70
6	-1.94	-2.38	-1.45

Results are averaged over the levels of: topo

Point estimate displayed: median

Results are given on the logit (not the response) scale.

HPD interval probability: 0.95

```
emmeans(df.seed.3, ~time, type="response")
```

time	response	lower.HPD	upper.HPD
1	0.00934	0.00291	0.0188
2	0.05570	0.03103	0.0873
3	0.07878	0.04880	0.1218
4	0.12035	0.07660	0.1713
5	0.10606	0.06341	0.1488
6	0.12567	0.07929	0.1803

Results are averaged over the levels of: topo

Point estimate displayed: median
Results are back-transformed from the logit scale
HPD interval probability: 0.95

We could look at pairwise difference between all levels of `time` (all campaigns), but `emmeans` also can display only consecutive contrasts with `consec~...`

```
emmeans(df.seed.3, consec~time)
```

```
$emmeans
  time emmean lower.HPD upper.HPD
1    -4.66   -5.57   -3.86
2    -2.83   -3.37   -2.30
3    -2.46   -2.94   -1.95
4    -1.99   -2.43   -1.53
5    -2.13   -2.60   -1.70
6    -1.94   -2.38   -1.45
```

Results are averaged over the levels of: topo
Point estimate displayed: median
Results are given on the logit (not the response) scale.
HPD interval probability: 0.95

```
$contrasts
  contrast      estimate lower.HPD upper.HPD
time2 - time1    1.840    0.902    2.812
time3 - time2    0.378   -0.287    1.045
time4 - time3    0.469   -0.132    1.147
time5 - time4   -0.150   -0.770    0.451
time6 - time5    0.190   -0.424    0.817
```

Results are averaged over the levels of: topo
Point estimate displayed: median
Results are given on the log odds ratio (not the response) scale.
HPD interval probability: 0.95

The contrasts between the last 3 campaigns are not different from 0 (95% CIs widely cover 0), so we can answer the research question with **yes**.

Alternatively, you could fit a model `Predation ~ seed + topo + time.combined`, where `time.combined` is the same as `time`, but levels 4,5,6 are combined in 1 level. Then do a model comparison versus the full model.